



Solution to Avoid Race Conditions

Thread Synchronization (Locking)



- Thread synchronization

- Prevents the **withdrawal thread** from withdrawing money when the **deposit thread** is depositing money.
- Prevents the **deposit thread** from depositing money when the **withdrawal thread** is withdrawing money.



- Thread synchronization

- Synchronizes threads

- Allows **only one thread** to access a shared variable at a time

- Forces all other threads to **wait and take turn** to access it.

- Enables **serialized** (or **mutually-exclusive**) access to a shared variable

- ```
public void deposit(double amount){
 balance = balance + amount; }
public void withdraw(double amount){
 balance = balance - amount; }
```

- Thread synchronization prevents race conditions by eliminating the following cases:

- Has read **balance**, but hasn't read **amount** yet.  
**Context switch** occurs, followed by a balance change by another thread
  - Has read **balance** and **amount**, but hasn't done addition/subtraction yet.  
**Context switch** occurs, followed by a balance change by another thread
  - Has read **balance** and **amount** and finished addition/subtraction, but hasn't updated **balance** yet  
**Context switch** occurs, followed by a balance change by another thread

# Thread Synchronization with Java

- Java implements **thread synchronization** by
  - Providing **locks**
  - Allowing you to write **atomic code** (a.k.a critical section) with locks.
    - **Atomic code**: A piece of code that can be executed by multiple threads in a **serialized** (or **mutually-exclusive**) manner.
      - Only one thread can run it at a time.
        - » When a thread is running it, all other threads have to wait to run it
  - Threads access a shared variable in atomic code.

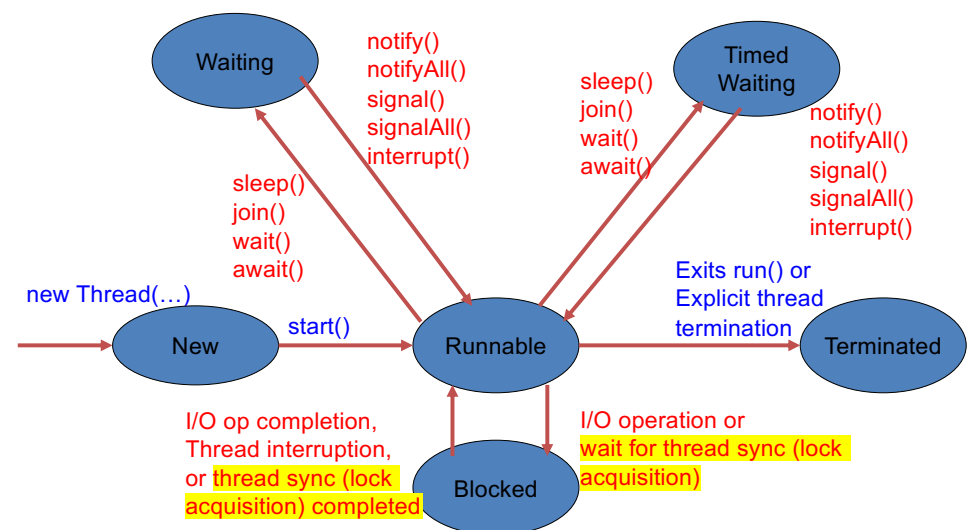
# Locks in Java

- Used to synchronize (or serialize, or mutually-exclude) multiple threads that access shared data.
- `java.util.concurrent.locks.Lock` interface
  - **ReentrantLock** class: the most commonly-used class for thread synchronization
    - Defines methods for threads to access shared data in a synchronized (or serialized, or mutually-exclusive) manner.
- Atomic code is surrounded by `lock()` and `unlock()` method calls on a lock.
  - ```
ReentrantLock aLock = new ReentrantLock();  
aLock.lock();  
atomic code (access to shared data)  
aLock.unlock();
```

6

- Once a thread calls `lock()` on a **lock**.
 - it acquires and holds the **lock** until it calls `unlock()`.
 - No other threads can acquire the lock until it is released with `unlock()`.
 - No other threads can run atomic code until the lock is released.
- If a thread calls `lock()` on a **lock** when another thread already holds the **lock**,
 - it goes to the **blocked** state and *gets blocked* (cannot do anything further) until the lock is released.

States of a Thread



7

8

How Can a Blocked Thread Run Again?

- JVM **automatically** let a *blocked thread* run again when the target lock is released.
 - You (your code) don't have to do anything for that.
- JVM's thread scheduler
 - Periodically reactivates all blocked threads so that they can try to acquire the target lock.
 - If the lock is still unavailable, they get blocked again.
 - Or, detects a release of the target lock.
 - May notify all blocked threads so that one of them can acquire the target lock.
 - May choose one of the blocked threads to acquire the lock.
 - Each blocked thread can eventually acquire the target lock.

9

Coding Idiom

- You MUST call `unlock()` to release a lock after running atomic code.
- Call `unlock()` in a **finally clause**.
 - ```
ReentrantLock aLock = new ReentrantLock();
aLock.lock();
try {
 atomic code (access to a shared variable)
}finally{
 aLock.unlock(); }
```
- If a `unlock()` call is NOT placed in a finally clause, it may NOT be invoked in the worst cases:
  - if `run()` returns in atomic code
  - if atomic code throws an exception

10

```
aLock.lock();
try{
 atomic code
}
finally{
 aLock.unlock();
}
```


**DO THIS!**

```
try{
 aLock.lock();
 atomic code
}
finally{
 aLock.unlock();
}
```

**DON'T DO THIS!**


- Make sure to call `lock()` **BEFORE** the "try" clause.
- If a thread throws an exception in `lock()`, it will not acquire the lock. However, it will call `unlock()`.
  - e.g., `lock()` can throw an `InterruptedException` when another thread call `interrupt()`.

## ThreadSafeBankAccount.java




- ```
private ReentrantLock lock = new ReentrantLock();
```
- ```
public void deposit(double amount){
 lock.lock();
 try{
 balance += amount; // atomic code
 }finally{
 lock.unlock(); } }
```
- ```
public void withdraw(double amount){
    lock.lock();
    try{
        balance -= amount; // atomic code
    }finally{
        lock.unlock(); } }
```

12



• Thread synchronization

- Prevents the **withdrawal thread** from withdrawing money when the **deposit thread** is depositing money.
- Prevents the **deposit thread** from depositing money when the **withdrawal thread** is withdrawing money.



- ```
public void deposit(double amount){
 balance = balance + amount;
}
```

  - A compound of 5 atomic operations.

## • Thread synchronization enables **serialized** (or **mutually-exclusive**) execution of a compound operation.

- Allows **only one thread** to perform the compound operation at a time.

```
- ReentrantLock aLock = new ReentrantLock();
aLock.lock();
try{
 balance = balance + amount; // atomic code
}
finally{
 aLock.unlock();
}
```

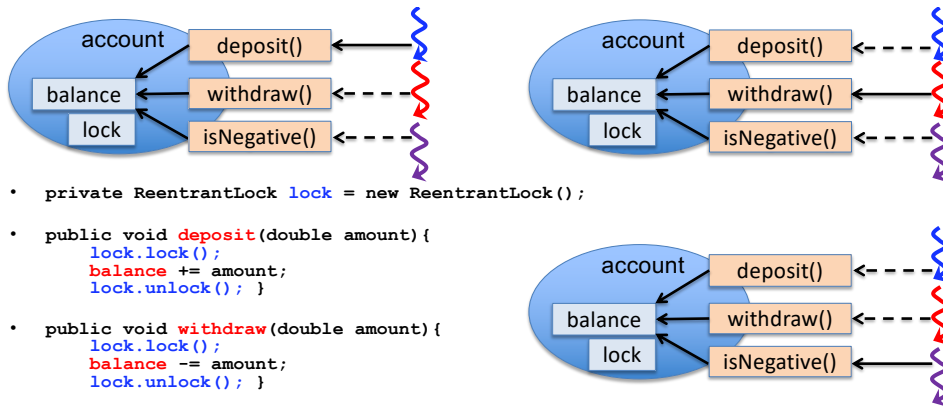
## Summary: How to Avoid Race Conditions?

- When multiple threads share and access a variable concurrently,
  - Make sure to **guard the shared variable with a lock**.
    - Identify **ALL read and write logic** to be performed on the variable
    - Surround each of them with `lock()` and `unlock()` to be called on the **same lock**

```
- e.g., public void deposit(double amount){
 lock.lock();
 try{
 balance = balance+amount; // atomic code (read & write)
 }finally{
 lock.unlock();
 }
}

- public void withdraw(double amount){
 lock.lock();
 try{
 balance = balance-amount; // atomic code (read & write)
 }finally{
 lock.unlock();
 }
}
```

- ```
public void deposit(double amount){
    ReentrantLock aLock = new ReentrantLock();
    aLock.lock();
    try{
        balance = balance + amount; // atomic code
    }finally{
        aLock.unlock();
    }
}
```
- As a result, the deposit and withdrawal threads update **balance** in either of the following 2 cases:
 - Has't read **balance**.
Context switch occurs, followed by a balance change by another thread
 - Has read **balance** and **amount**, performed addition/sutraction, and updated **balance**
Context switch occurs, followed by a balance change by another thread



- ```
private ReentrantLock lock = new ReentrantLock();
```
- ```
public void deposit(double amount) {
    lock.lock();
    balance += amount;
    lock.unlock();
}
```
- ```
public void withdraw(double amount) {
 lock.lock();
 balance -= amount;
 lock.unlock();
}
```
- ```
public boolean isNegative() {
    lock.lock();
    if (balance < 0) { return true; }
    else if { return false; }
    lock.unlock();
}
```
- **Thread synchronization**
 - Mutually excludes those 3 threads, so only one of them can run atomic code associated with “lock” at a time.
 - It is important to use the **same lock** in **ALL read and write logic** to be performed on “balance.” Otherwise, threads are NOT mutually excluded.

Nested Locking

- ```
class BankAccount {
 private double balance;
 private double MIN_BALANCE = ...;
 private double PENALTY = ...;
 private ReentrantLock lock = ...;

 public void withdraw(double amount) {
 lock.lock();
 balance -= amount; // 5+ steps
 if (balance < MIN_BALANCE) // 3+ steps
 subtractPenaltyFee();
 lock.unlock();
 }

 private void subtractPenaltyFee() {
 balance -= PENALTY; // 5+ steps

 // NO NEED TO SURROUND THIS LINE with LOCK() and UNLOCK()
 // because it is called from atomic code.
 }
}
```

18

- ```
class BankAccount {
    private double balance;
    private double MIN_BALANCE = ...;
    private double PENALTY = ...;
    private ReentrantLock lock = ...;

    public void withdraw(double amount) {
        lock.lock();
        balance -= amount;           // 5+ steps
        if (balance < MIN_BALANCE)    // 3+ steps
            subtractPenaltyFee();
        lock.unlock();
    }

    private void subtractPenaltyFee() {
        balance -= computePenalty(); // many steps

        // NO NEED TO SURROUND THIS LINE with LOCK() and UNLOCK()
        // because it is called from atomic code.
    }

    private void computePenalty() {
        ...                          // many steps

        // NO NEED TO SURROUND THIS LINE with LOCK() and UNLOCK()
        // because it is called from atomic code.
    }
}
```

19

Thread Reentrancy

- ```
class BankAccount {
 private double balance;
 private double MIN_BALANCE = ...;
 private double PENALTY = ...;
 private ReentrantLock lock = ...;

 public void withdraw(double amount) {
 lock.lock();
 balance -= amount; // 5+ steps
 if (balance < MIN_BALANCE) // 3+ steps
 subtractPenaltyFee();
 lock.unlock();
 }

 private void subtractPenaltyFee() {
 lock.lock();
 balance -= PENALTY; // 5+ steps
 lock.unlock();
 }
}
```

- This code does not have a (deadlock) problem.
- A thread can **re-enter** (or **re-acquire**) the same lock as far as it already owns the lock.

20

```

class A{
 private B b = new B();
 private ReentrantLock lock=...;

 public void a1(){
 lock.lock();
 b.b1(this); //nested locking
 lock.unlock();
 }
 public void a2(){
 lock.lock();
 do something.
 lock.unlock();
 }
}

```

```

Class B{
 public void b1(A a){
 a.a2();
 }
}

```

If a thread performs:  
 A a = new A();  
 a.a1();  
 it *re-enters (or re-acquires)* the same lock that it already owns.

- This code does not have a (deadlock) problem.
- A thread can *re-enter (or re-acquire)* the same lock as far as it already owns the lock.

## RunnableCancellablePrimeGenerator

```

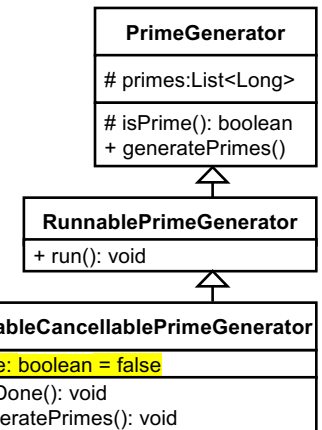
class RunnableCancellablePrimeGenerator ...{
 private boolean done = false;

 public void setDone(){
 done = true;
 }

 public void generatePrimes(){
 for(long n = from; n <= to; n++){
 if(done){
 System.out.println("Stopped...");
 this.primes.clear();
 break;
 }
 if(isPrime(n)){ this.primes.add(n); }
 }
 }

 public void run(){
 generatePrimes();
 }
}

```



### Main thread

### Thread t

```

gen = new RunnableCancellablePrimeGenerator(...)
t = new Thread(gen)

t.start()

gen.setDone()

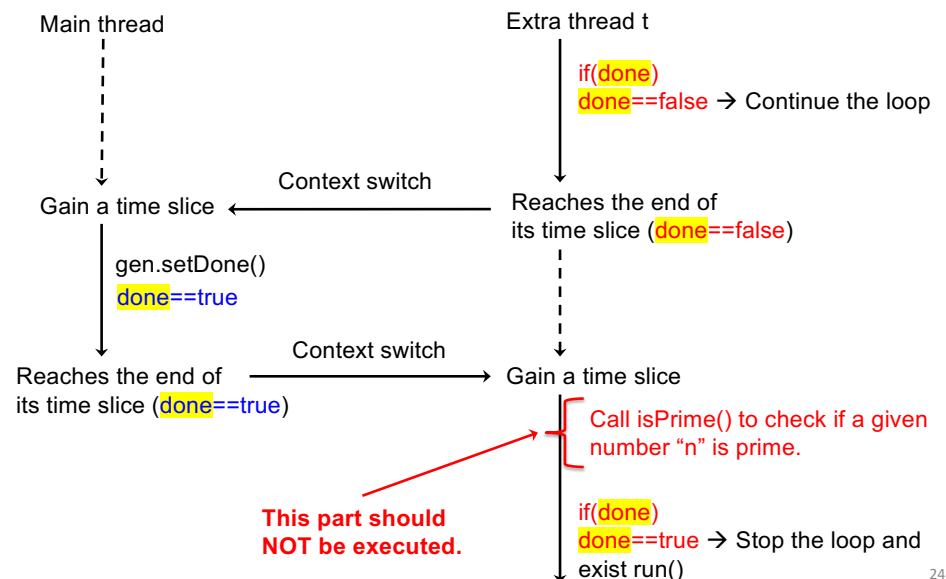
for(long n = from; n <= to; n++){
 if(done){
 System.out.println("Stopped generating prime nums.");
 this.primes.clear();
 break;
 }
 if(isPrime(n)){ this.primes.add(n); }
}

```

Thread t: Executes gen.run() to generate prime nums  
 Prints "stopped generating prime nums" and exits run()

- This code is **NOT thread-safe**. Race conditions can occur.

## A Potential Race Condition



```

class RunnableCancellablePrimeGenerator extends PrimeNumberGenerator{
 private boolean done = false;

 public void run(){
 for(long n = from; n <= to; n++){
 if(done){
 System.out.println("Stopped generating prime nums.");
 this.primes.clear();
 break;
 }
 if(isPrime(n)){ this.primes.add(n); }
 }
 }

 public void setDone(){
 done = true;
 }
}

```

Thread t  
Context switch

25

```

class CancellablePrimeNumberGenerator extends PrimeNumberGenerator{
 private boolean done = false;

 public void run(){
 for(long n = from; n <= to; n++){
 if(done){
 System.out.println("Stopped generating prime nums.");
 this.primes.clear();
 break;
 }
 if(isPrime(n)){ this.primes.add(n); }
 }
 }

 public void setDone(){
 done = true;
 }
}

```

Thread t  
Context switch

Main thread  
Context switch

26

## Visibility Issue

```

class CancellablePrimeNumberGenerator extends PrimeNumberGenerator{
 private boolean done = false;

 public void run(){
 for(long n = from; n <= to; n++){
 if(done){
 System.out.println("Stopped generating prime nums.");
 this.primes.clear();
 break;
 }
 if(isPrime(n)){ this.primes.add(n); }
 }
 }

 public void setDone(){
 done = true;
 }
}

```

Thread t  
Context switch

Main thread  
Context switch

27

- The most up-to-date value of the shared variable “done” is not visible for all threads.
- Solution:
  - Identify all read and write logic on the shared variable
  - Surround each read/write logic with lock() and unlock() calls on the same **ReentrantLock**

# Risk of Race Conditions

- Race conditions generate different risk levels in different applications.
  - It may be acceptable to make an extra loop iteration and generate an extra prime number.
  - It is NOT acceptable to make an extra loop iteration in other applications in
    - Safety/mission critical domains
      - Nuclear plant control, autonomous driving, rocket launch, weapon management, etc.
    - Financial domains
      - Banking, high-frequency trading, stock exchange management, etc.
    - Medical/healthcare domains
- In any case, you should know that a race condition exists and should be able to fix it.

# How to Prevent Potential Race Conditions

- When multiple threads share and access a variable concurrently.
  - Make sure to guard the shared variable
    - By surrounding each read/write logic on it with lock() and unlock().
- When a loop checks a condition with a shared variable (e.g., flag).
  - Surround the conditional with lock() and unlock()
  - Try NOT to surround the entire loop with with lock() and unlock()! Why?
    - Does not enjoy concurrency.
    - May result in a deadlock.

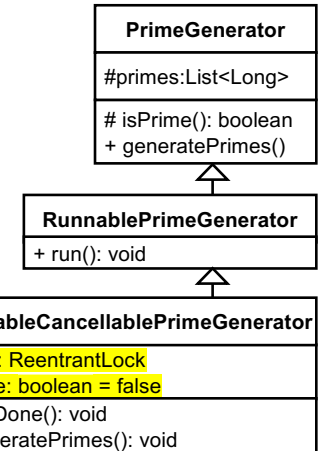
# Solution w/ Thread Synchronization

```
class RunnableCancellablePrimeGenerator ...{
 private ReentrantLock lock = new ReentrantLock();
 private boolean done = false;

 public void setDone(){
 lock.lock();
 try{
 done = true;
 }finally{
 lock.unlock();
 }
 }

 public void generatePrimes(){
 for(long n = from; n <= to; n++){
 lock.lock();
 try{
 if(done) break;
 if(isPrime(n)) {this.primes.add(n); }
 }finally{
 lock.unlock();
 }
 }
 }

 public void run(){
 generatePrimes();
 }
}
```



# Having the Entire Loop as Atomic Code Does NOT Enjoy Concurrency

- DO **NOT** do this.

```
lock.lock();
try{
 for(n = from; n <= to; n++){
 if(done) break;
 if(isPrime(n)) {
 this.primes.add(n);
 }
 }
}finally{
 lock.unlock();
}
```
- Do this.

```
for(n = from; n <= to; n++){
 lock.lock();
 try{
 if(done) break;
 if(isPrime(n)) {
 this.primes.add(n);
 }
 }finally{
 lock.unlock();
 }
}
```
- ```
lock.lock();
try{
    done = true;
}finally{
    lock.unlock();
}
```


- If a thread acquires the lock to start generating primes, it will **release that lock after all primes are generated.**
 - No other threads can flip the flag until all primes are generated.
 - No other threads can terminate the prime generation thread when it is generating primes.

A thread that generates primes

```
lock.lock();
try{
    for(n = from; n <= to; n++){
        if(done) break;
        if(isPrime(n)){
            this.primes.add(n);
        }
    }finally{
        lock.unlock();
    }
}
```

The main thread, which wants to terminate the prime generation thread, can acquire the lock after all primes have been generated.

```
lock.lock();
try{
    done = true;
}finally{
    lock.unlock();
}
```

- If a thread acquires the lock to start printing #s, it will print #s forever.
 - No other threads can flip the flag forever (deadlock!)

```
lock.lock();
while(!done){
    System.out.println("#"); // read logic
}
lock.unlock();
```

The purple thread gets stuck here forever because the green thread never release the lock.

```
lock.lock();
done = true; // write logic
lock.unlock();
```

Having the Entire Loop as Atomic Code May Result in a Deadlock

- DO **NOT** do this.

```
lock.lock();
try{
    while(!done){
        System.out.println("#");
    }
}finally{
    lock.unlock();
}
```

```
lock.lock();
try{
    done = true;
}finally{
    lock.unlock();
}
```

- Do this.

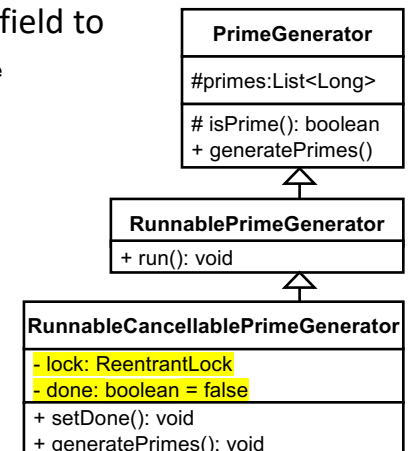
```
while(true){
    lock.lock();
    try{
        if(done) break;
        System.out.println("#");
    }finally{
        lock.unlock();
    }
}
```

```
lock.lock();
try{
    done = true;
}finally{
    lock.unlock();
}
```

HW 6-1

- Revise `RunnableCancellablePrimeGenerator` to be thread-safe.

- Add a `ReentrantLock` as a data field to guard the shared variable `done`
- Revise `generatePrimes()` and `setDone()` to access `done` with the lock

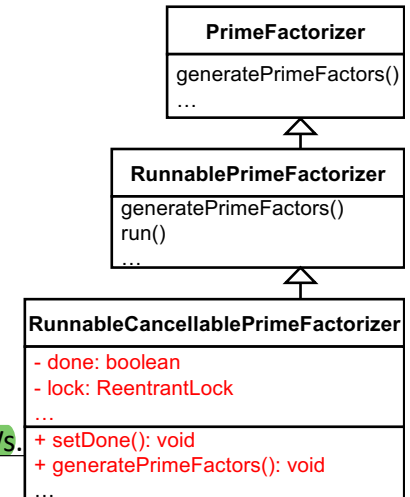


HW 6-2

- In `setDone()`,
 - Use try-finally blocks.
 - Call `unlock()` in a finally block. Always do this in all subsequent HWs.
- In `generatePrimes()`
 - Use try-finally blocks; Call `unlock()` in a finally block
 - Do not surround the entire “for” loop with `lock()` and `unlock()`.

- Implement `RunnableCancellablePrimeFactorizer` by extending `RunnablePrimeFactorizer`.

- Implement flag-based thread termination
- Add a flag variable `done`
- Add a `ReentrantLock` to guard the shared variable `done`
- Revise `generatePrimeFactors()` and `setDone()` to access `done` with the lock
- Call `unlock()` in a finally block. Always do this in all subsequent HWs.




Recap: Singleton Design Pattern

- Place your HW6-1 and HW6-2 solutions in the same package: `edu.umb.cs681.primes`.
- Upload them to the same GitHub repo.
- Guarantee that a class has only one instance.
 - c.f. CS680 lecture note
- ```

public class Singleton{
 private Singleton(){};
 private static Singleton instance = null;

 // Factory method to return the singleton instance
 public static Singleton getInstance(){
 if(instance==null)
 instance = new Singleton();
 return instance;
 }
}

```
- This code is NOT thread-safe; race conditions can occur.




- When multiple threads call `getInstance()` concurrently, they share the data field `instance`.

```
public class Singleton{
 private Singleton(){};
 private static Singleton instance = null;

 // Factory method to return the singleton instance
 public static Singleton getInstance(){
 if(instance==null) // Read: (multiple) steps
 instance = new Singleton(); // Write: (multiple) steps
 return instance; // Read: (multiple) steps
 } // NOT THREAD-SAFE
}
```

41



- ```
public class Singleton{
    private Singleton(){};
    private static Singleton instance = null;
                                // Write logic. Requires 2 steps
                                // THREAD-SAFE }
}
```
- JVM completes all **initial value assignments on all static data fields** BEFORE using a class or creating class instances.
 - `instance` has been initialized before any threads call `getInstance()`
 - This write logic is thread-safe.
 - No need to worry about race conditions here.

42

Concurrent Singleton Design Pattern

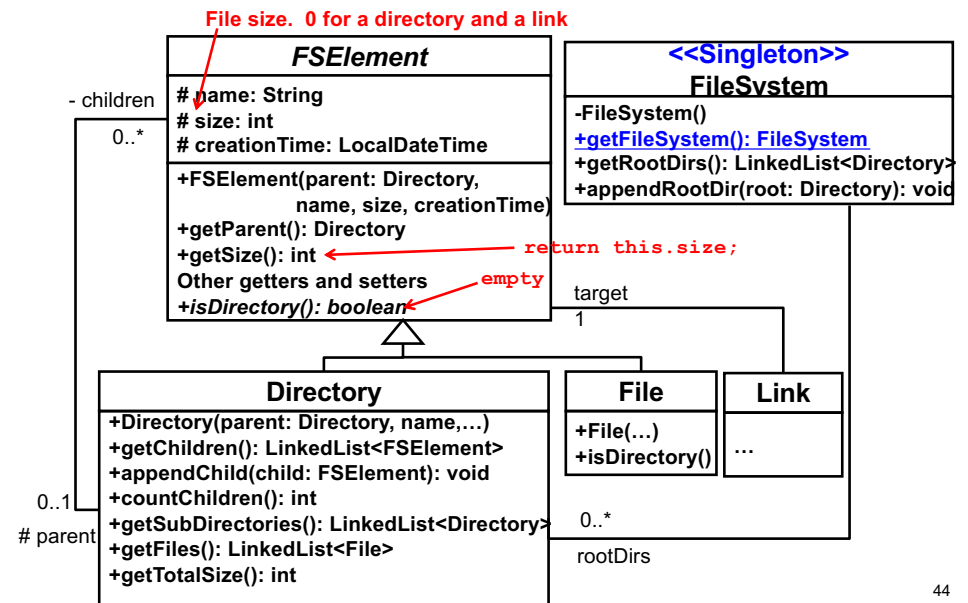
- Guarantee that a class has only one instance.

```
public class ConcurrentSingleton{
    private ConcurrentSingleton(){};
    private static ConcurrentSingleton instance = null;
    private static ReentrantLock lock = new ReentrantLock();

    // Factory method to create or return the singleton instance
    public static Singleton getInstance(){
        lock.lock();
        try{
            if(instance==null){ instance = new ConcurrentSingleton(); }
            return instance;
        }finally{
            lock.unlock();
        }
    }
}
```

43


Recall CS680's HW 8



44

Concurrency in a File System

- Modern file systems are all multi-threaded.
 - Any modern OSes want apps and OS services to access a file system concurrently.
 - e.g., reading/updating metadata of file system elements and adding/removing file system elements
 - e.g., saving a file while renaming another file.



HW 7

- Revise `FileSystem.getInstance()` to be thread-safe
 - Define a lock in `FileSystem`. Use the lock in `getInstance()` to guard the `instance` in data field.
 - Use try-finally blocks: Always do this in all subsequent HWs.
 - Run multiple threads to call `getInstance()`
 - Make sure that only one instance is created.
 - c.f. You worked on identify check in CS680.
- No need to revise other methods of `FileSystem`.
 - e.g., `getRootDirs()`, `appendRootDir()`, etc.


Let's Make your File System Multi-threaded

- In CS680's HW 8, you wrote basic classes to implement a single-threaded file system
 - `FSElement`, `Directory`, `File`, `Link` and `FileSystem`
 - You never considered multiple threads.
- To make your file system multi-threaded, you need to expect multiple threads will use them concurrently.
 - For example, multiple threads will use `FileSystem` concurrently.
 - e.g., They will call `FileSystem.getInstance()` concurrently.
 - But, `FileSystem` is **NOT thread-safe**.


Just in Case... Static Data Fields and Methods

- A **static** data field
 - Created and used on a **per-class** basis.
 - All instances share the data field.
- A **regular (non-static)** data field
 - Created and used on an **instance-by-instance** basis.
 - Different instances have different copies of the data field.

- A **static** data field
 - Created and used on a **per-class** basis.
 - All instances share the data field.
- A **regular (non-static)** data field
 - Created and used on an **instance-by-instance** basis.
 - Different instances have different copies of the data field.



- A **static** method
 - Created and used on a **per-class** basis.
 - Can access static data fields.
 - **Can NOT** access regular (non-static) data fields.
- A **regular (non-static)** method
 - Created and used on an **instance-by-instance** basis.
 - Can access **both** regular (non-static) and static data fields.



Just in case... Initialization of Data Fields

- ```

• public class Foo{
 ReentrantLock lock = new ReentrantLock();
 static ReentrantLock sLock = new ReentrantLock();
 int foo = 10;
 int boo;
 Object obj; }

```
- Data fields (both static and non-static ones) are initialized **before** a constructor is called.
    - Data field initialization is always **thread-safe**.
    - The default value is used for initialization if the initial value is not explicitly given.
      - e.g., 0 for int, null for a reference type, etc.