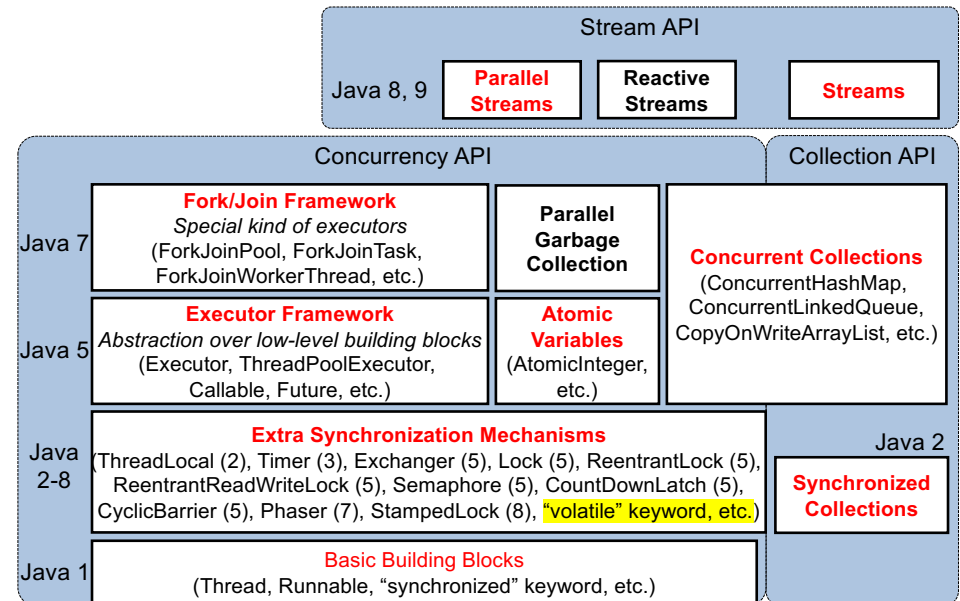


Volatile Variables

Concurrency API in Java



Recap: Thread Synchronization (Locking)

- You must do thread synchronization (locking) to guard a variable shared among threads.
 - Otherwise, race conditions can occur.

What is the “volatile” Keyword?

- You can **skip thread synchronization** to guard a shared variable **in some cases** by defining the variable as **volatile**.
 - No race conditions occur on the variable even though you never do thread sync.
 - Greater peace of mind regarding thread safety
 - Less coding/maintenance effort
 - Less overhead (higher performance)
 - Thread sync consumes some CPU time and memory space.

Recap: A Race Condition in Explicit Thread Termination

```
boolean done = false;

public void setDone(){
    done = true;
}

public void run(){
    while( true )
        if( done ) break;
    else{
        // Do some work
    }
}
```

- A thread (“A”: soon-to-be-terminated)
 - Has read the current value of `done`, which is `false` (in blue), but...
 - Has NOT applied the value (false) to a conditional (in yellow) yet
- A context switch occurs.
- Another thread (“B”: terminator) calls `setDone()` to assign `true` to `done`.

- Another context switch occurs.
- The value change that “B” just made is NOT visible for “A”.
 - “A” applies the value of `done` (false) to the conditional (in yellow) and runs green code instead of breaking out from the loop.

Read on a Volatile Variable

```
volatile boolean done = false;

public void setDone(){
    done = true;
}

public void run(){
    while( true )
        if( done ) break;
    else{
        // Do some work
    }
}
```

- A context switch can still occur in b/w
 - Reading the value of `done` (in blue) and
 - Applying it to the conditional (in yellow).
- Then, Thread “B” (terminator thread) may call `setDone()` to assign `true` to `done`.
- After another context switch, Thread “A” re-loads the most up-to-date value of `done` from the memory
 - because it’s a volatile variable.

- A volatile variable makes a value change on a shared variable visible for all threads.

Thread-safe Code

- With thread sync

```
boolean done = false;
ReentrantLock lock = new ReentrantLock();

public void setDone(){
    lock.lock();
    done = true;
    lock.unlock();
}

public void run(){
    while(true){
        lock.lock();
        try{
            if( done ) break;
            // Do some work
        }finally{
            lock.unlock();
        }
    }
}
```

- With a volatile variable

```
volatile boolean done = false;

public void setDone(){
    done = true;
}

public void run(){
    while( true )
        if( done ) break;
    // Do some work
}
```

- Both versions are thread-safe.
- A volatile variable makes code simpler and less error-prone.

6

Write on a Volatile Variable

```
volatile boolean done = false;

public void setDone(){
    done = true;
}

public void run(){
    while( true )
        if( done ) break;
    // Do some work
}
```

- A “terminator” thread #1 called `setDone()`
 - Has loaded the value of true, but...
 - Has NOT assigned it to `done` yet.
- Then, a context switch occurs.
- Another “terminator” thread #2 calls `setDone()`
 - Completes an assignment of true to `done`
- After another context switch, thread #1 assigns true to `done`.
- No race conditions occur.
 - Because two threads (#1 & #2) assign the same value.

Special Effect on a Volatile Variable

- A volatile variable is **guaranteed to have the most up-to-date value** whenever it is read/used.

```
- volatile int a;
int b = a;           // These 2-step ops are all thread-safe,
println(a);          // even if a context switch occurs in between
                      // the 2 steps and another thread changes
                      // the value of "a" there.

if(a==0)              // 3 steps. It is thread-safe, even if a
                      // context switch occurs in b/w the 1st and 2nd
                      // steps, or in b/w the 2nd and 3rd steps, and
                      // another thread changes the value of "a"
                      // there.

a + 1;               // 3 steps: reading the value of "a", loading 1
                      // and summing up "a" and 1.
                      // This 3-step op is thread-safe, even if a
                      // context switch occurs in b/w the 1st and 2nd
                      // steps, or in b/w the 2nd and 3rd steps, and
                      // another thread changes the value of "a"
                      // there.
```

- This “special effect” is called **“memory semantics”** or **“memory effect”** of a volatile variable in Java API doc and other resources.

```
- long a;
a = 10L;              // Requires 2 atomic steps because long is
                      // a 64-bit type. NOT thread-safe.
                      // A context switch can occur during the 10L
                      // value assignment. The first 32-bit half of
                      // value may be assigned by a thread, and the
                      // second half may be assigned by another
                      // thread.

- volatile long a;
a = 10L;              // Thread-safe even if a context switch occurs
                      // during the 10L value assignment.
```

- This “special effect” is called **“memory semantics”** or **“memory effect”** of a volatile variable in Java API doc and other resources.

10

Limited Effectiveness

- A volatile variable does **NOT** eliminate all possible race conditions...

```
- volatile int a;
a + 1;                // 3 steps. Thread-safe.

b = a + 1;            // These operations are NOT thread-safe.
if(a+1>0)              // The first 3 steps (in yellow) are
println(a+1);          // thread-safe though.
                      // The 3rd step generates an intermediate
                      // state that is not volatile. It may NOT be
                      // in synch with the most up-to-date value of
                      // "a", if a context switch occurs in b/w the
                      // 3rd and 4th steps and another thread changes
                      // the value of "a" there.
```

- A volatile variable is effective **only when no intermediate states/values are generated with it.**

- A volatile variable is effective **only when no intermediate states/values are generated with it.**

```
- volatile int a;
a = 1;                // 2 steps. Thread-safe.
                      // A context switch can occur in b/w the
                      // 2 steps, but no race conditions
                      // occur here.

a = a + 1;            // 5 steps. Not thread-safe.
a++;                  // The first 3 steps (in yellow) are
                      // thread-safe though.
                      // The 3rd step generates an intermediate
                      // state that is not volatile. It may NOT be
                      // in synch with the most up-to-date value of
                      // "a", if a context switch occurs in b/w the
                      // 3rd and 4th steps and another thread changes
                      // the value of "a" there.
```

- The “volatile” keyword imposes a **strong/strict limitation** in its use cases.
 - It works only for a variable whose value is used without any intermediate state.
 - Good: if(a>0), if(done), a=1
 - NOT good: if(a+1>0), if(!done), b=a+1, a=a+1
- Carefully use it only when you can live with this limitation.

When to Use Volatile Variables?

- So, what use cases generate no intermediate states/values?
 - Probably, **latch** only.
 - Performs **a single type of value changes**
 - e.g. False → True, Non-zero → Zero, Negative → Non-negative, null to non-null, etc.
 - It is used in explicit thread termination.
 - The **flag variable done** in prior examples
 - » The state of **done** always changes in a unidirectional way: false → true.
 - » “true → false” never happen.
 - The flag **done** is a perfect example to be volatile effectively.

- The keyword “volatile” works for **all primitive types**
 - Both 32-bit and 64-bit types
- However, it DOES NOT work for **arrays and reference types**.

In Summary

- The “volatile” keyword is **NOT a general-purpose, widely-applicable** tool.
- Powerful only **in some specific cases**
 - In practice, assume it is useful only for implementing a latch in a thread-safe manner.
 - **Very useful** to implement **flag-based thread termination** and **2-step thread termination** in a thread-safe manner.
- Note that there exist numerous projects that misuse “volatile” variables.

Exercise (Not HW)

- Declare a flag as a volatile variable in flag-based and 2-step thread termination schemes that you have implemented.
 - e.g., prime generation and factorization
 - As far as the flag is “volatile,” you don’t have to guard it with a lock.

Note: Do NOT Use “volatile” to Implement Concurrent Singleton


- Not thread-safe

```
public class Singleton{
    private Singleton(){};
    private static Singleton instance = null;

    public static Singleton getInstance(){
        if(instance==null)
            instance = new Singleton();
        return instance;
    }
}
```

- Can we turn the *Singleton* class to be thread-safe by changing the `instance` variable to be `volatile`?
 - Its value changes unidirectionally: null to non-null.
- No, unfortunately,
 - Because it’s a `reference-type` variable.

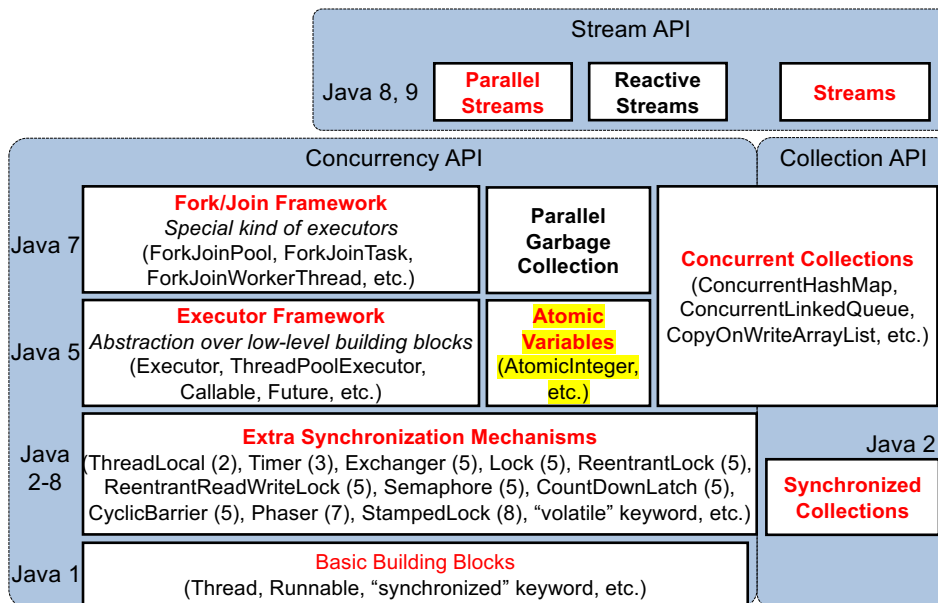
Note: Do NOT use “volatile” for a Variable that allows Bi-directional Value Changes



- File
 - Has a boolean variable: `changed`
 - Initialized to be `false`.
 - `change()`
 - Changes the file’s content.
 - Assigns `true` to `changed`.
 - `save()`
 - If `changed==false`, return;
 - If `changed==true`
 - Save the file’s content.
 - Assign `false` to `changed`.

Atomic Data Structures

Concurrency API in Java



java.util.concurrent.atomic Package

- Offers **thread-safe classes** to manipulate single variables.
 - AtomicBoolean,
 - AtomicInteger, AtomicIntegerArray
 - AtomicLong, AtomicLongArray
 - AtomicReference<V>, AtomicReferenceArray<E>
 - DoubleAccumulator, DoubleAdder
 - LongAccumulator, LongAdder
 - ...

java.util.concurrent.atomic Package

- Offers **thread-safe classes** to manipulate single variables.
 - AtomicBoolean,
 - AtomicInteger, AtomicIntegerArray
 - AtomicLong, AtomicLongArray
 - AtomicReference<V>, AtomicReferenceArray<E>
 - DoubleAccumulator, DoubleAdder
 - LongAccumulator, LongAdder
 - ...
- All of their public methods are implemented to be **thread-safe**. They **avoid race conditions**.

Atomic Variables

- Serve as **"better" volatile variables**.
- Offer the **"special" memory effects as volatile variables do**.
- Provide additional support for **atomic value changes**.
 - get()** has the memory effect of reading a volatile variable.
 - Returned value is guaranteed to be **the most up-to-date** whenever it is used.
 - set()** has the memory effect of writing to a volatile variable.
 - Value assignment is atomic.
 - read-and-update** methods (e.g., **xxxAndSet()** and **getAndxxx()**) have the memory effects of reading and writing to volatile variables.
- Highly recommended as far as they match your use cases.

AtomicBoolean

- Offers the methods to manipulate a single boolean value atomically (i.e., in a thread-safe manner).

```
- boolean flag = false;
  if(flag)
    println(flag);
    // 2 steps. Not thread-safe
    // 2 steps. Not thread-safe
```

```
- volatile boolean vFlag = false;
  if(vFlag)
    println(vFlag);
    // 2 steps. Thread-safe
    // 2 steps. Thread-safe
```

```
- AtomicBoolean atomicFlag = new AtomicBoolean(false);
  if(atomicFlag.get())
    println(atomicFlag.get());
    // 2 steps. Thread-safe
    // 2 steps. Thread-safe
```

- `get()`: Thread-safe. Atomically returns the current value, with the memory effect of reading a volatile variable.
 - A context switch can occur right after `get()` returns, and another thread may change the boolean value.
 - However, `AtomicBoolean` guarantees that the most up-to-date value will be re-loaded, whenever it is used.

```
- volatile boolean vFlag = false;
  vFlag = true;
  if(vFlag){ int i=10; }
  // 2 steps. Thread-safe.
  // Thread-safe.

  if(vFlag){ vFlag = false; } // NOT thread-safe.
  // Read and write ops are thread-safe
  // each, but a race condition can occur
  // if a context switch happens in b/w
  // the two ops (in yellow and blue) and
  // vFlag is updated.
```

```
- AtomicBoolean atomicFlag = new AtomicBoolean(false);
  atomicFlag.set(true)
  atomicFlag.compareAndSet(true, false);
  // Thread-safe
  // Thread-safe!
```

- `set()`: Thread-safe. Atomically sets a value.

— *read-and-update* methods (e.g., `xxxAndSet()` and `getAndXxx()`) have the memory effects of both reading and writing volatile variables.

AtomicBoolean for Thread Termination

- `boolean compareAndSet(boolean expect, boolean update)`

— Atomically sets the “update” (new) value if the current value is equal to the “expect” value.

— Returns true if successful.

— Returns false if the current value was not equal to the “expect” value.

- `atomicFlag.compareAndSet(true, false);` // Thread-safe
 - Sets false if the current value is true, and returns true
 - Keeps the current value if it is false, and returns false

- `AtomicBoolean` can be used to implement a flag variable in flag-based and 2-step thread termination schemes

```
class CancelableRunnable
implements Runnable {
```

```
    boolean done = false;
    ReentrantLock lock;
```

```
    public void setDone(){
        lock.lock();
        done = true;
        lock.unlock();
    }
```

```
    public void run(){
        while(true){
            lock.lock();
            if(done) break;
            lock.unlock();
            ... // do some work }}}
    }
```

```
class CancelableRunnableWithAtomicBoolean
implements Runnable {
```

```
    AtomicBoolean done =
        new AtomicBoolean(false);
```

```
    public void setDone(){
        done.set(true);
    }
```

```
    public void run(){
        while(true){
            if(done.get()) break;
            ... // do some work
        } }
    }
```

AtomicInteger

- Use `AtomicBoolean` or “volatile” to implement explicit thread termination in all future HWs.
- Exercise: Use `AtomicBoolean` in explicit thread termination in prior sample programs

- Offers thread-safe methods to manipulate a single integer value atomically.

```
- int i = 0;
  int j = i;
  if(i==0)                                     // 2 steps. Not thread-safe.
                                              // 3 steps. Not thread-safe.
```

```
- volatile int i = 0;
  int j = i;
  if(i==0)                                     // 2 steps. Thread-safe.
                                              // 3 steps. Thread-safe.
```

```
- AtomicInteger atomicInt = new AtomicInteger(0);
  int j = atomicInt.get();                    // Thread-safe. j==0
  if(atomicInt.get()==0){...}                // Thread-safe.
```

- `get()`: Thread-safe. Atomically returns the current value, with the memory effect of reading a volatile variable.

- A context switch can occur right after `get()` returns, and another thread may change the value.
- However, `AtomicInteger` guarantees that the most up-to-date value will be re-loaded, whenever it is used.

```
- volatile int i = 0;
  i = 1;                                     // Thread-safe.
  int j = i + 1;                             // 5 steps. NOT thread-safe.
  i++;                                       // 5 steps. NOT thread-safe.

- AtomicInteger atomicInt = new AtomicInteger(0);
  atomicInt.set(1);                          // Thread-safe.
  j = atomicInt.incrementAndGet();            // Thread-safe. j==2
  atomicInt.incrementAndGet();                // Thread-safe.
```

- `set()`: Thread-safe. Atomically sets a value., with the memory effect of writing to a volatile variable.

- *read-and-update* methods (e.g., `xxxAndSet()`, `xxxAndGet()` and `getAndXxx()`) have the memory effects of both reading and writing volatile variables.

```
- volatile int i = 0;
  i = 1;                                     // Thread-safe.

  if(i==0){i = 1;}                          // NOT thread-safe.
                                              // Read and write ops are thread-safe
                                              // each, but a race condition can occur
                                              // if a context switch happens in b/w
                                              // the two ops (in yellow and blue) and
                                              // the variable is updated.
```

```
- AtomicInteger atomicInt = new AtomicInteger(0);
  atomicInt.compareAndSet(0, 1);             // Thread-safe.
```

- Other *read-and-update* methods include:
 - `decrementAndGet()`, `addAndGet(int)`, `getAndSet(int)`, ...
 - `updateAndGet(IntUnaryOperator)`, `accumulateAndGet(int, IntBinaryOperator)`

– `updateAndGet(IntUnaryOperator updateFunction)`

- *Atomically* updates the current integer value with the result of running a given LE and returns the updated value.
- `IntUnaryOperator`: a general-purpose functional interface
- `updateFunction`: a LE that takes the current `int` value as a parameter, updates it and returns the updated value.
 - This update logic runs atomically
- ```
AtomicInteger atomicInt = new AtomicInteger(10);
atomicInt.updateAndGet((int i)-> ++i); // 11. Thread safe
atomicInt.incrementAndGet(); // 12. Thread safe
```

|                                     | Params           | Returns          |
|-------------------------------------|------------------|------------------|
| <code>UnaryOperator&lt;T&gt;</code> | <code>T</code>   | <code>T</code>   |
| <code>IntUnaryOperator</code>       | <code>int</code> | <code>int</code> |

- ```
volatile i = 10;
if(i > 0){ i = 1;}
```

// NOT thread-safe.
 // Read and write ops are thread-safe
 // each, but a race condition can occur
 // if a context switch happens in b/w
 // the two ops (in yellow and blue) and
 // the variable is updated.
- ```
AtomicInteger atomicInt = new AtomicInteger(10);
atomicInt.updateAndGet((int i)->{ if(i > 0){ i = 0; }
 return i; });
```

// Thread safe. The lambda expression  
 // is executed in a thread-safe manner.

- ```
AtomicInteger atomicInt = new AtomicInteger(10);
atomicInt.updateAndGet( (int i)-> ++i ); // 11. Thread safe
```

• **Why ++1?** Just in case, note that:

```
- int i = 0;
  i++;           // i==1

- int i = 0;
  int x = 0;
  int x = i++;   // i==1, x==0

- int i = 0;
  int y = 0;
  int y = ++i;   // i==1, y==1
```

– `accumulateAndGet(int, IntBinaryOperator)`

- *Atomically* updates the current `int` value with the result of running a given LE and returns the updated value.
- `IntBinaryOperator`: a general-purpose functional interface
- ```
atomicInt.accumulateAndGet(initValue,
 (result, currentVal)-> ...);
```
- ```
int result = initValue;
result = accumulate(result, currentVal);
```
- This LE runs atomically (i.e., in a thread-safe manner)

	Params	Returns
<code>BinaryOperator<T></code>	<code>T, T</code>	<code>T</code>
<code>IntBinaryOperator</code>	<code>int, int</code>	<code>int</code>

AtomicReference<V>

```
• AtomicInteger atomicInt = new AtomicInteger(0);
  atomicInt.accumulateAndGet(10,
    (result, currentVal)->
      (currentVal+result)/2 );

    // Returns (0+10)/2

• AtomicInteger atomicInt = new AtomicInteger(0);
  atomicInt.accumulateAndGet(
    0, (result, currentVal)->{
      if(currentVal >= result) return currentVal;
      else if return result; }));
```

- Offers thread-safe methods to manipulate a **reference type value** atomically.

– Note: “volatile” does NOT work for reference types.

```
– Car car = new Car(...);
  Car anotherCar = car;           // 2 steps. Not thread-safe.
  if(car!=null)                   // 3 steps. Not thread-safe.

– AtomicReference<Car> atomicCar =
  new AtomicReference<>(new Car(...)); // Thread-safe
  atomicCar.set(new Car(...));         // Thread-safe
  if(atomicCar.get() != null)          // Thread-safe
```

– Other useful methods

- compareAndSet(V,V), updateAndGet(UnaryOperator<V>), accumulateAndGet(V, BinaryOperator<V>), etc.
- C.f. API documentation

Recap: Concurrent Singleton

- Guarantee that a class has only one instance.
 - Its static factory method is implemented to be thread-safe because it performs thread sync.

```
• public class ConcurrentSingleton{
  private ConcurrentSingleton(){};
  private static ConcurrentSingleton instance = null;
  private static ReentrantLock lock = new ReentrantLock();

  // Factory method to create or return the singleton instance
  public static Singleton getInstance(){
    lock.lock();
    try{
      if(instance==null){ instance = new ConcurrentSingleton(); }
      return instance;
    }finally{
      lock.unlock();
    }
  }
}
```

Exercise (Not HW)

- Revise ConcurrentSingleton
 - Use AtomicReference to define instance
 - Take out ReentrantLock. No need to do thread sync!
 - Implement getInstance() in a thread-safe manner with AtomicReference. Use compareAndSet() or updateAndGet()
- Test it:

```
public static void main(String[] args){
  for(int i=0; i<10; i++){
    new Thread(
      ()->{Sys.out.println(ConcurrentSingleton.getInstance());}).start();
  }
}
```