

Immutable Classes

Concurrency and Immutability

- Classes that **never change the state of each instance**
 - State = A set of data field values
 - Getter methods only; **no setter methods** available.
- No state (value) changes → **No race conditions!**
 - Race condition = State/value inconsistency on a shared data field(s) among threads
 - C.f. `BankAccount`'s `balance`, `PrimeGenerator`'s `done`
- No methods require thread sync.
- All methods are always thread-safe.

2

An Example Immutable Class:

`java.lang.String`

- ```
char[] str = {"u", "m", "b"};
String string = new String(str);
```
- ```
String string = "umb";    // Syntactic sugar for the above code
```
- A series of constructors to initialize string data.
- All non-constructor methods **never change the initialized string**.
- **No setter methods** are available.
 - c.f. API doc.

Example Methods in String

- ```
String str = "umb";
System.out.println(str); // println() calls str.toString()
 // prints out "umb"
```
- ```
System.out.println( str.replace("b","l")); // uml
                               // Creates a new String instance that
                               // contains "uml" and returns it.
```
- ```
System.out.println(str); // umb
```
- ```
System.out.println( str.toUpperCase() ); // UMB
                               // Creates a new String instance
                               // that contains "UMB" and returns it.
```
- ```
System.out.println(str); // umb
```
- ```
System.out.println( str.substring(0,1) ); // um
                               // Creates a new String instance that contains
                               // "um" and returns it.
```
- ```
System.out.println(str); // umb
```
- Some methods of `string` look like setters, but they are actually NOT.
  - They never change the initialized string data ("umb").

3

# String

- **Final class**, which cannot be extended (sub-classed)
  - `public final class String{...}`
  - Prevents its sub-classes from introducing new methods that try to change the initialized string data (e.g. “umb”).
- Maintains the initialized string in a **private and final data field**.
  - `public final class String{  
    private final char value[];  
    ... }`
  - Once a value is assigned to a final variable, the value cannot change afterward.
    - No methods of `String` can change the initialized string data.

- Each “setter-like” method of `String` creates another `String` instance that contains another string data.

```
- public final class String{
 private final char[] value; // Immutable
 ...
 public String toUpperCase(){
 int length = value.length; // Local variable
 char[] result = new char[length]; // Local variable

 for(int i = 0; i < length; i++){ // Local variable
 result[i] = ... // Transform value[i] to an upper case
 }
 return new String(result); // Local variable
 } }
```

- This is actually NOT a setter method! It is thread-safe.
  - **No local variables are shared by threads.**
    - Each thread creates a copy of a local variable.
    - Different threads use different copies of it.

## Benefits of Immutability

- For API designers
  - An immutable class never require thread sync in its methods.
    - No need to guard its data field (e.g., `value` in `String`) with a lock
      - The data field’s value never changes.
      - All threads simply read “fixed” (or “finalized”) data from the data field.
  - Its methods are **free from race conditions**.
    - Code revision and debugging get easier.
- For API users
  - Immutable classes are free **from race conditions**.
    - Their client code doesn’t always need thread synchronisation.
  - They are **free from potential performance loss** due to thread sync.
    - Thread sync forces every thread to acquire a lock.
      - There is some overhead to acquire a lock.
    - If the lock is not available, the thread needs to get into the Blocked state until it becomes available.
      - It cannot do anything to make progress.

## Note That...

- An immutable class's methods are thread-safe, but...
- Client code of those methods may NOT be always thread-safe.
  - The code below is **NOT thread-safe**; it requires thread sync.

```
public class Person {
 private String firstName, lastName; // Shared variables

 public void setLastName(String last){
 lastName = last; } // 2 Steps; NOT thread-safe

 public String getLastName(){
 return lastName; } // 2 Steps; NOT thread-safe

 public String getFullName(){
 return firstName + " " + lastName; // Multi-steps; NOT thread-safe
 }
 // Syntactic sugar for new StringBuilder().append(firstName).append(" ")
 // .append(lastName).toString()
 // Reads on firstName and lastName are NOT thread-safe.
 // An instance of StringBuilder is local for each thread though.
```

- An immutable class's methods are thread-safe, but...
- Client code of those methods may NOT be always thread-safe.
  - The code below **is thread-safe**.

```
public class ErrorMsgGenerator {
 ...
 public String getFileNotFoundErrMsg(String path){
 return "The requested file " + path + " is not found.";

 // Syntactic sugar for new StringBuilder().append("The requested...")
 // .append(path)
 // .append(" was not...")
 // .toString();
 // Multiple steps, but thread-safe
 // Local variables are not sharable by threads.
 // Reads and writes on a local variable are thread-safe.
 // "path" is a local variable.
 // An instance of StringBuilder is local for each thread.
 } }
```

## Other Immutable Classes

- Wrapper classes for primitive types
  - c.f. CS680
- `java.nio.file.Path`
  - c.f. CS680
- `java.util.regex.Pattern`
- Some classes in `java.net`
  - e.g., `URL`, `URI`, `Inet4Address` and `Inet6Address`
- Date and Time API (`java.time`)

| Primitive type | Wrapper class |
|----------------|---------------|
| boolean        | Boolean       |
| byte           | Byte          |
| char           | Character     |
| float          | Float         |
| int            | Integer       |
| long           | Long          |
| short          | Short         |
| double         | Double        |

## Integer

- Wrapper class of an `int` value
  - **Final class**, which cannot be extended (sub-classed)
  - Maintains the initialized `int` data in a **private and final data field**.
    - `Integer int = Integer.valueOf(10);` // Static factory method
    - `Integer int = 10;` // Auto-boxing; Syntactic sugar for the above code
  - Has no setter methods; **no methods change the initialized `int` data**.
  - **All methods are thread-safe**.

## Note That...

- An immutable class's methods are thread-safe, but...
- Client code of those methods may NOT always be thread-safe.
  - The code below is NOT thread-safe; it requires thread sync.

```
public class Person {
 private Integer age; // Shared variable

 public void setAge(Integer age){
 this.age = age; } // 2 steps; NOT thread-safe

 public Integer getAge(){
 return this.age; } // 2 steps; NOT thread-safe

 public boolean isKindergartener(){ // Multi-steps. Not thread-safe
 if(this.age < 6){return true;} // Equivalent to:
 else{return false;} // if(this.age.intValue()<6)...
 }
```

## Implementing User-Defined (Your Own) Immutable Classes

- Immutable class
  - Defined as a **final class**
  - Has **private final data fields** only.
  - Has **no setter** methods.
- Clearly state immutability in program comments, API documents, design documents, etc.
  - Java API documentation does so too.
  - Put {immutable} to a class in UML

- An immutable class's methods are thread-safe, but...
- Client code of those methods may NOT always be thread-safe.
  - The code below is thread-safe.

```
public class ErrorMsgGenerator {
 private final Integer FILE_NOT_FOUND = Integer.valueOf(404);
 ...

 public String getFileNotFoundErrorMsg(Path path){
 String header = "Error code: " + FILE_NOT_FOUND;
 // Syntactic sugar for
 // new StringBuilder().append("Error code: ")
 // .append(FILE_NOT_FOUND.toString()).toString()
 // Reads on local and final variables are thread-safe.
 // An instance of StringBuilder is local for each thread.

 String body = "The requested file " +
 path.toString() + " is not found."
 // "path" is a local variable. A read on it is thread-safe.

 return header + " " + body; // Thread-safe } }
```

## An Example User-Defined Immutable Class

```
public final class SSN {
 private final int first3Digits, middle2Digits, last4Digits;

 public SSN(int first, int middle, int last){ // Thread-safe
 this.first3Digits = first;
 this.middle2Digits = middle;
 this.last4Digits = last; }

 public int getLast4Digits(){return last4Digits;} // 2 steps but a read
 // on a final variable
 // is thread-safe.

 public String toString(){
 return first3Digits + "-" + middle2Digits + "-" + last4Digits;
 // Multiple steps, but thread-safe
 // Reads on final variables are thread-safe. }

 public boolean equals(SSN anotherSSN){
 if(this.toString().equals(anotherSSN.toString())){ return true; }
 else{ return false; }
 // Multiple steps, but thread-safe
 // SSN and String are immutable. } }
```

## Note That...

- `public final class SSN {`  
    `private final int first3Digits, middle2Digits, last4Digits;`  
  
    `public SSN(int first, int middle, int last){ // Thread-safe`  
        `this.first3Digits = first;`  
        `this.middle2Digits = middle;`  
        `this.last4Digits = last; }`
- A constructor is always executed **atomically**.
  - Only one thread can run a constructor on a class instance that is being created and initialized.
    - Multiple threads never call a constructor(s) on the same instance concurrently.
  - Until a thread returns/completes a constructor on a class instance, no other threads can call public methods on that instance.

- An immutable class's methods are thread-safe, but...
- Client code of those methods may not always be thread-safe.
  - The code below is **NOT thread-safe**; it requires thread sync.

```
public class Person {
 private SSN ssn; // Shared (non-final) variable

 public Person(SSN ssn){ this.ssn = ssn; }

 public SSN setSSN(SSN ssn){
 this.ssn = ssn; // 2 steps; not thread-safe
 }

 public SSN getSSN(){ // 2 steps; not thread-safe
 return ssn; // "ssn" is NOT final.
 }
}
```

Person requires thread sync to guard `ssn`, although `SSN` does not.

- An immutable class's methods are thread-safe, but...
- Client code of those methods may NOT always be thread-safe.
  - The code below **is thread-safe**.

```
public class Person {
 private final SSN ssn; // Shared final variable

 public Person(SSN ssn){ this.ssn = ssn; }

 public SSN getSSN(){ // 2 Steps, but thread-safe.
 return ssn; // "ssn" is final.
 }
}

Person person = new Person(new SSN(012, 34, 5678));
person.getSSN();
```

## Record Type: Convenience Mechanism to Implement an Immutable Classes

- Available since Java 16

```
public record SSN(int first3Digits,
 int middle2Digits,
 int last4Digits){}
```

- Syntactic sugar for:

```
public final class SSN extends java.lang.Record{
 private final int first3Digits;
 private final int middle2Digits;
 private final int last4Digits;

 public SSN(int first3Digits, int middle2Digits, int last4Digits){
 this.first3Digits = first3Digits;
 ... }

 public int first3Digits(){ return first3Digits; }

 // 2 other getters, equals(), toString() and hashCode() are
 // auto-generated }
```

# HW 9

- You can
  - Override the (auto-generated) constructor.
  - Override auto-generated methods (e.g. getters)
  - Add extra methods.

```
• public record SSN(int first3Digits,
 int middle2Digits,
 int last4Digits){
 public int to9Digits(){
 return ...
 }
}
```

```
public class Aircraft {
 private Position position; // Shared (non-final) variable

 public Aircraft(Position pos){ this.position = pos; }

 public void setPosition(double newLat,
 double newLong,
 double new Alt){
 this.position = this.position.change(newLat, newLong, newAlt);
 // 2 steps; NOT thread-safe. }

 public Position getPosition(){
 return position; // 2 steps; NOT thread-safe.
 } }

Aircraft requires thread sync to guard position, although
Position does not.
```

- Implement your own immutable:

```
public record Position(double latitude,
 double longitude,
 double altitude){} // decimal degrees
```

- Add extra methods:

```
List<Double> coordinate()
```

```
Position change(double newLat, double newLon, double newAlt){
 return new Position(newLat, newLon, newAlt);}
```

- It sounds like a setter, but it is NOT. It creates a new instance and returns it.

```
boolean higherAltThan(Position anotherPosition)
boolean lowerAltThan(Position anotherPosition)
boolean northOf(Position anotherPosition)
boolean southOf(Position anotherPosition)
```

- Turn in

- Immutable Position

- Thread-safe Aircraft, which...

- has a **ReentrantLock** as a data field and performs thread synch with it, OR
- uses **AtomicReferenceType<Position>** and skips thread sync.

- Runnable class whose run() calls Aircraft's setPosition() and getPosition()

- You can replace the Runnable class with a lambda expression, if you like.

- Test code to create and run multiple threads

# Performance Implications

- An immutable object makes a big difference in performance
  - when threads often **read** data from the object
  - because they do not perform thread synch.
- If you are interested, compare the performance of
  - Immutable `Position` and
  - Mutable `Position` that performs thread sync in its setters and getters.
  - Immutable `Position` is approx. 25% faster on my machine.
- An immutable object never trigger performance loss in single-threaded apps.
  - If a single-threaded app calls a mutable object's method that performs thread synch, the app incurs unnecessary performance loss.
    - The app never need thread synch, but the mutable object's method does it for the app.
- Recent APIs often rely on immutables, whenever/wherever possible, considering both thread safety and performance.
  - e.g., Date and Time API
- Setter-like methods of an immutable class
  - e.g., `replace()` and `toUpperCase()` of `String`
  - NOT that efficient
    - because they create new instances of the immutable class.
- Again, the performance benefit of an immutable class stands out when...
  - threads often **READ** data from an immutable object

## Well, Not All Classes can be Immutable...

- Immutable classes are good for both API designers and API users.
- However, in practice, many classes need to be mutable...
- Think of **separating a class to mutable and immutable parts**
  - if **read operations** are called **very often**
  - if **write operations** are **occasionally** called but they are **computationally expensive**.

## An Example: String and StringBuilder

- Both represent string data.
- **String**
  - **Immutable**: Its state never change once it is set up.
  - Thread-safe
  - **Getter methods** (reads): **thread-safe**.
  - **Setter-like methods** (writes): **thread-safe**, **relatively slow**.
- **StringBuilder**
  - **Mutable**: Its state can change through its methods.
  - NOT thread-safe; its public methods never perform thread synch.
  - **Getters** (reads): **NOT thread-safe**.
  - **Setters** (writes): **NOT thread-safe**, **relatively fast**.

29

## Performance Comparison: String Concatenation as an Example Write Operation

- ```
String str = "UMass";
str = str + " Boston";           // "UMass Boston"
// Syntax sugar for:
// str = new StringBuilder(str).append(" Boston").toString();

// Creates 2 instances: StringBuilder and String
// Calls 3 methods: StringBuilder's constructor and 2 methods
```
- ```
StringBuilder builder = new StringBuilder(str);
builder.append(" Boston");
str = builder.toString();
```
- No difference in performance.

- Use **String** as far as...
  - you often call getter methods (i.e. read operations)
    - e.g., `toString()`, `equals()`, `contains()` and `length()`
  - you occasionally call setter-like methods (i.e. write operations)
    - e.g., `replace()`, `toUpperCase()` and `substring()`
    - Here, “occasionally” means calling a setter-like method once or a few times at a time.
- Use **StringBuilder** only when...
  - you call write operations more often than usual
    - e.g., hundreds/thousands of times in a loop
    - **StringBuilder's** write operations (setter methods) are faster than **String's** write operations (setter-like methods).

- ```
String header = "Error code: " + FILE_NOT_FOUND;
String body = "The requested file " + path.toString()
              + " was not found."

return header + " " + body;
// Syntax sugar for:
// header = new StringBuilder("").append(FILE_NOT_FOUND).toString();
// body = new StringBuilder("").append(...).append(" ").toString();
// return new StringBuilder(header).append(" ").append(body).toString();
//
// Creates 6 instances and calls 11 methods
```
- ```
StringBuilder builder = new StringBuilder();
builder.append("Error code: ");
builder.append(FILE_NOT_FOUND);
builder.append("The requested file ");
...
builder.append(" was not found.");
return builder.toString();
// Creates 2 instance and calls 5 methods
```

- More visible difference in performance, if string concatenation is performed with **multiple** statements.



- `LinkedList<String> emailAddrs = ...;`
- `String` `commaSeparatedEmailAddrs;`  
`for(String emailAddr: emailAddrs){`  
`commaSeparatedEmailAddrs += emailAddr + ", "; }`
- `StringBuilder` `commaSeparatedEmailAddrs;`  
`for(String emailAddr: emailAddrs){`  
`commaSeparatedEmailAddrs.append(emailAddr).append(", "); }`
- The latter code can run **20-100% faster** depending on the number of collection elements (i.e. email addresses).
- **DO NOT use `string` (immutable class) to perform a number of write operations.**

## StringBuffer

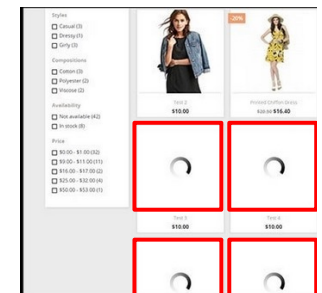
- Provides the same set of public methods as `StringBuilder` does.
- `StringBuffer` (since Java 1.0)
  - All public methods are **thread-safe** with thread sync.
  - Client code of `StringBuffer` may still require thread sync.
  - DO NOT use this class.
    - It makes no sense to use it in single-threaded apps.
- `StringBuilder` (since Java 5)
  - All public methods are **NOT thread-safe**.
  - Client code of `StringBuilder` require thread sync.
  - Use this class
    - regardless of single-threaded or multi-threaded apps.

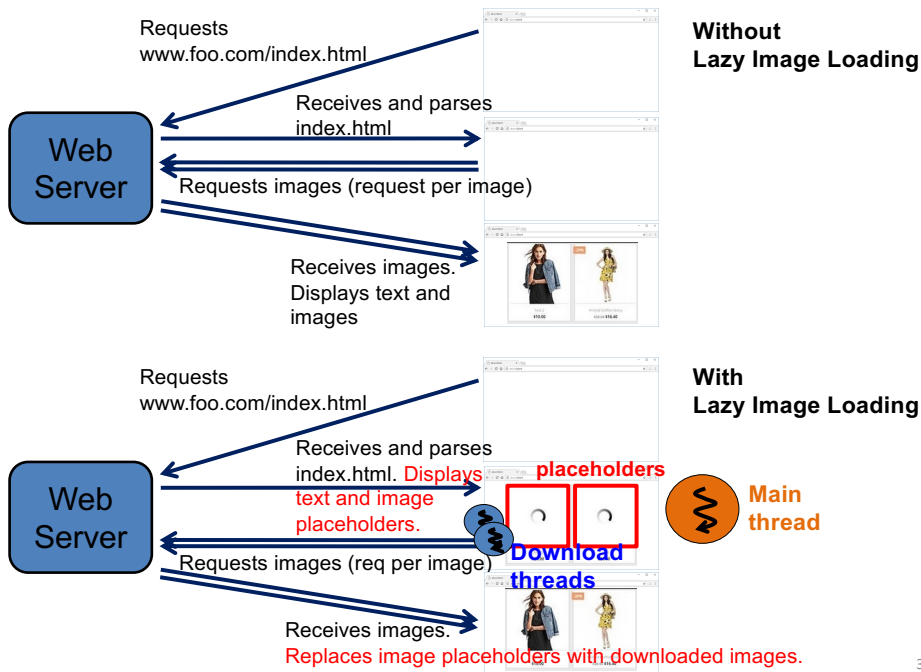
## When to Use Immutable and Mutable Classes?

- Use `string` (immutable class) for **read** operations
- Use `StringBuilder` (mutable class) for **write** operations
  - Note that `StringBuilder`'s methods are NOT thread-safe; e.g., `append()`.
    - Use a `StringBuilder` as a local variable
    - Do thread synchronization in multi-threaded apps.
- `String`-to-`StringBuilder` conversion is implemented in a constructor of `StringBuilder`.
- `StringBuilder`-to-`String` conversion is implemented in a constructor of `String`.

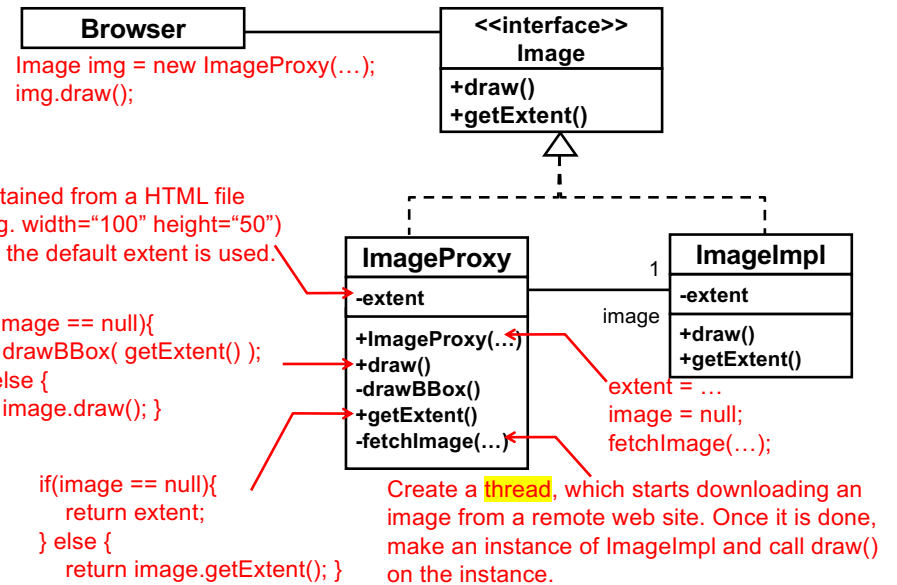
## Exercise: Lazy Image Loading

- Suppose a browser just downloaded an HTML file and found it contains images. The browser...
  - Displays a **bounding box (placeholder)** first for each image
    - Until it fully downloads the image.
      - Most users cannot be patient enough to keep watching blank browser windows until all text and images are downloaded and displayed.
  - Replaces the **bounding box** with the real image.

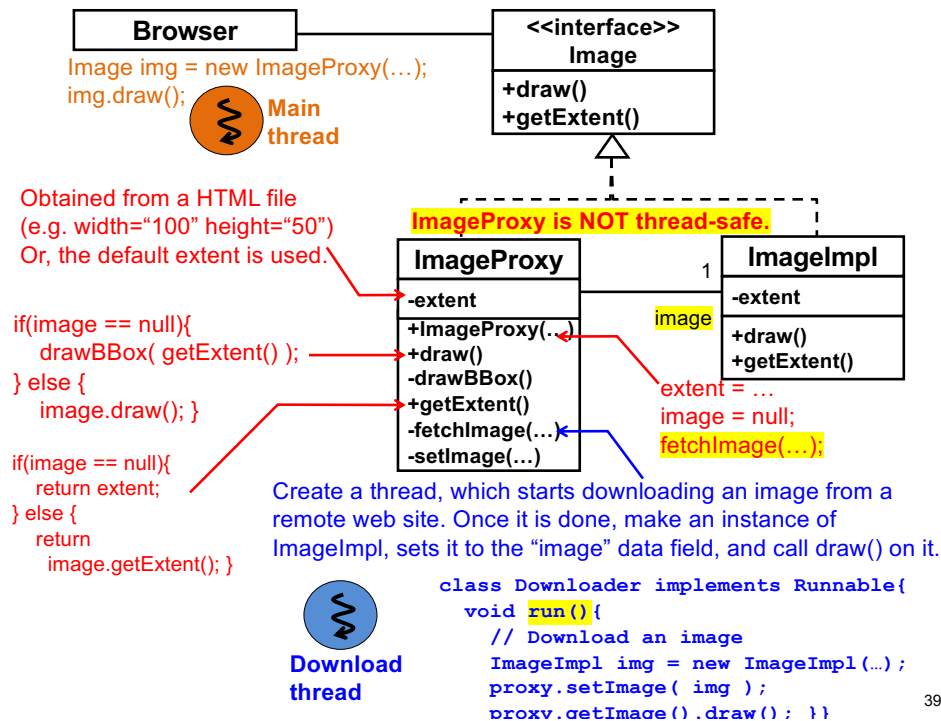




37



38



39

## Recap: Face Detection

- Suppose you are implementing an app to organize, edit and analyze pictures.
  - When the app loads a raw picture, it **superimpose a rectangle on a human face** in the picture by (dynamically) calling an external face detection/recognition API.
    - e.g., Microsoft Azure Face API, Google Cloud Vision API



- Some **delay** is expected to receive a face detection result from an external API.

- The user is not patient enough to keep watching a blank app window until the picture and face detection result are displayed.

- Lazy loading** of detection results

- Show the user a raw picture first.
- Call a face detection API.
- Receive a detection result.
- Replace the raw picture with a superimposed one, which contains a detection result.




**Client code (app):**  
 Picture pic = new RawPicture(...);  
 pic.draw();



path = ...  
 detectFaces(...);

if(superimposed == null){  
 drawRawImage(...);  
 } else {  
 superimposed.draw();  
 }

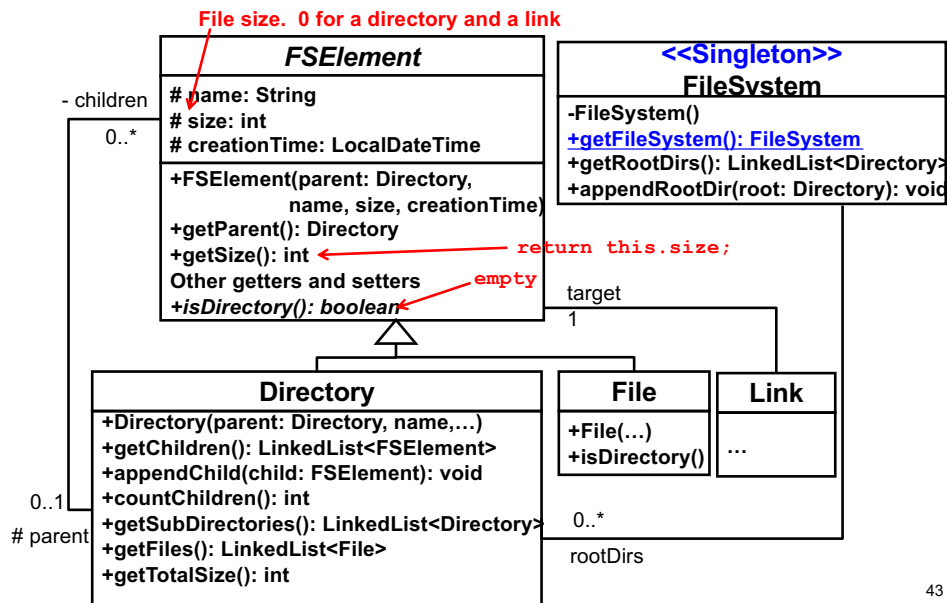


Create a thread, which calls an external API. Instantiate SuperimposedPicture and Face once the API returns a detection result. Set the SuperimposedPicture instance to the "superimposed" data field and call draw() on the instance.

**RawPicture is NOT thread-safe.**

42

## Recall CS680's HW 8



43

## HW 10

- In HW 7, you revised `FileSystem`'s `getInstance()` to be thread-safe.
- In this HW, revise all the methods in `FSElement`, `Directory`, `File`, `Link` and `FileSystem` to be thread-safe.
- Define a `ReentrantLock` in `FSElement` as its data field.
  - Use the lock to guard all the data fields in `FSElement`, `Directory`, `File`, and `Link`.
    - Any methods that access those data fields must perform thread sync with that lock.

– e.g., `name` in `FSElement`

```

getName() {
 this.lock.lock();
 return this.name;
 this.lock.unlock();
}

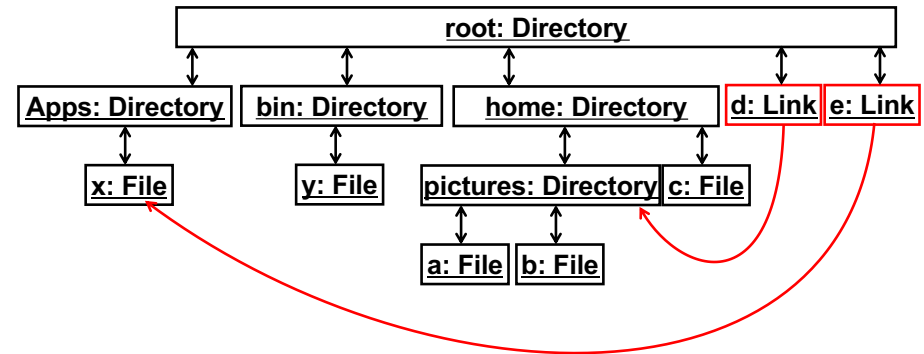
```

```

setName(String newName) {
 this.lock.lock();
 this.name = newName;
 this.lock.unlock();
}

```

- `creationTime` is typed with `LocalDateTime`.
  - All the classes in Date and Time API **are thread-safe**.
    - `LocalDateTime` is thread-safe; all of its public methods are thread-safe.
  - When you call a method of `LocalDateTime`, you do NOT have to do thread sync.
- `children` of `Directory` is typed with `LinkedList`.
  - All the collection classes in the `java.util` package are **NOT thread-safe**.
    - `LinkedList` is **NOT thread-safe**.
      - All of its public methods **are NOT thread-safe**; they never perform thread synchronization.
  - Whenever you call a method of `LinkedList` on children, make sure to do thread sync (i.e., surround that method call with `lock()` and `unlock()` on the lock in `Directory`).
- Keep `FileSystem` as a thread-safe *Singleton* class.



- Use this tree structure for testing.
  - You used this structure as a test fixture in CS680.
  - Assign values to data fields (size, etc) as you like.
- Have the main thread create and run 10+ extra threads, each of which will call the methods of **Directory**, **File**, **Link** and **FileSystem**.
- Have the main thread perform **2-step termination** for those extra threads.