

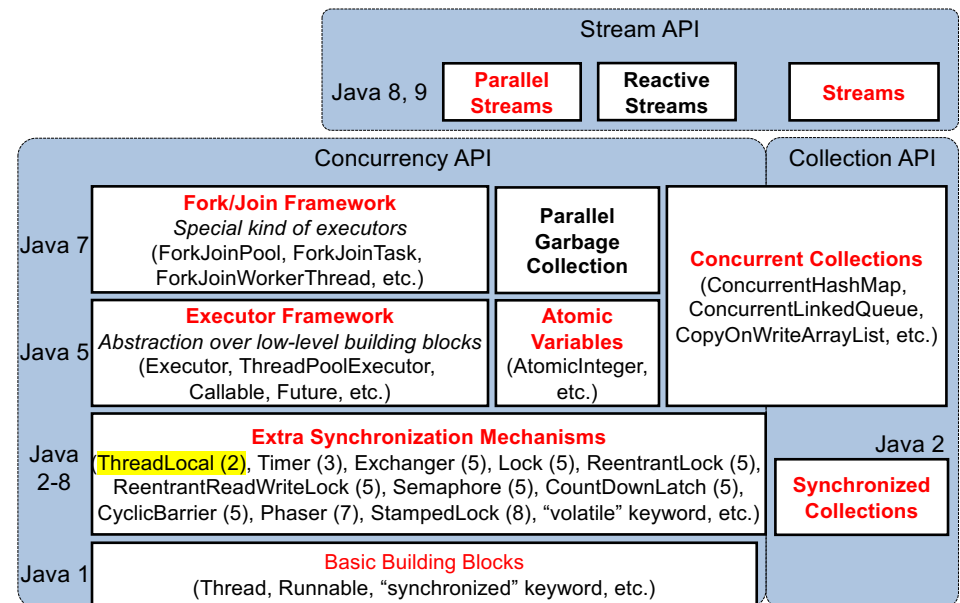
Thread-Specific Storage (TSS)

Thread-Specific Storage (TSS)


- Storage (memory space) that is allocated and reserved for a particular thread.
- Once a TSS is allocated to a thread, no other threads can access it.
 - It is dedicated for a single thread.
 - It is NEVER shared among multiple threads.
 - You **don't have to worry about potential race conditions** on it because of no sharing.
- Implemented in `java.lang.ThreadLocal<T>`

3

Concurrency API in Java




Common Use Case of ThreadLocal<T>



- Different threads generate different data of the same type (τ).
- Each of them
 - stores its own data into `ThreadLocal<T>` with `set()`
 - read its own data from the `ThreadLocal<T>` with `get()`


4



- **ThreadLocal**
 - allows threads to access data through their (thread) IDs.
 - A thread cannot access any data generated and maintained by the other threads.
 - Internally performs thread sync to guard the table.
- Client code `ThreadLocal` does not have to do thread sync. 5

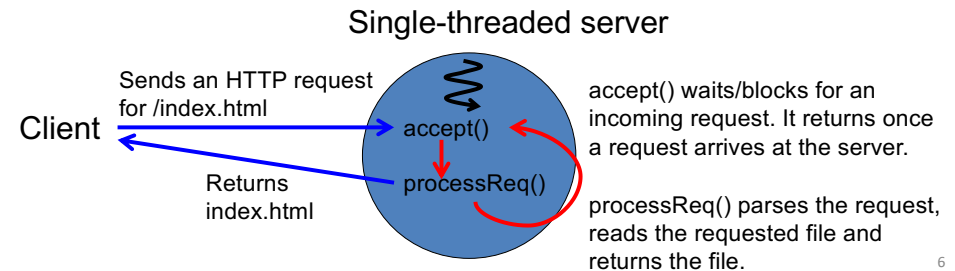
Thread-per-Request Concurrency

- Once the web server's main thread receives a request from a client, it creates a new thread.
 - The new thread parses the incoming request, reads the requested file and returns it.
 - The thread terminates once the requested file is returned to the client.





Exercise: Web Server Development


- Suppose you are developing a web server.
 - Receives an HTTP request that a client (browser) transmits to request an HTML file.
 - Returns the requested file to the client.
- What if the server receives multiple requests from multiple clients simultaneously?
 - If the server is single-threaded, it processes requests *sequentially*.



Concurrent (Multi-threaded) Web Server



TSS in a Concurrent Web Server



- Each request handler (thread):
 - May need **customer info** (e.g. **customer ID**) from a browser cookie to display some personalized content (e.g. shopping cart items)
 - May need **client-specific information** (e.g., **client OS name and browser name**) to display some client-specific content.

9

Race Conditions: A Summary

10

3 Cases where Race Conditions can Occur

- **CASE 1:** Threads **share and access a variable (data field)**
 - Solution 1: Define the variable as a **local** variable, rather than a data field.
 - Local variables are never shared among threads.
 - Solution 2: Define a **lock** as a data field of the variable's enclosing class and use it to access the variable
 - Surround every read/write logic on the variable with the `lock()` and `unlock()` calls on the lock
 - Solution 3: Use **thread-safe tools**
 - e.g. volatile variable, atomic variable, immutable class, TSS, etc.
 - Can skip thread synch
- **CASE 2:** Threads **call an API method that is NOT thread-safe.**
 - You cannot define a lock in the method's enclosing class (i.e., API class).
 - You need **"client-side locking,"** which means performing thread sync (or "locking") **in your client code** that uses the API method.

11

12

- Example: `LinkedList` is NOT thread-safe; its methods never perform thread sync.

```
- // This class is NOT thread-safe
public class LinkHolder {
    private LinkedList<Integer> list;

    public void addValue(Integer i){
        list.add(i);
    }

    public LinkedList<Integer> getList(){
        return list; } }
```

13

- Example: `LinkedList` is NOT thread-safe; its methods never perform thread sync.

```
- // This class is NOT thread-safe
public class LinkHolder {
    private LinkedList<Integer> list;

    public void addValue(Integer i){
        list.add(i);
    }

    public LinkedList<Integer> getList(){
        return list; } }

- // This class is thread-safe, thanks to client-side locking.
public class ListHolder {
    private LinkedList<Integer> list;
    private ReentrantLock lock = new ReentrantLock();

    public void addValue(Integer i){
        lock.lock();
        list.add(i);
        lock.unlock();
    }

    public LinkedList<Integer> getList(){
        lock.lock();
        return list;
        lock.unlock(); } }
```

14

- **CASE 3:** Threads run client code of a thread-safe method.

- You may (or may not) need “client-side locking,” which means performing thread sync in the client code.
- Example: `Integer` is thread-safe, but the following client code is not.

```
- // This class is NOT thread-safe.
public class Person {
    private Integer age;

    public void setAge(Integer age){
        this.age = age; }

    public Integer getAge(){
        return this.age; }

    public boolean isKindergartener(){ // Multi-steps. Not thread-safe
        if(this.age < 6){return true;} // Equivalent to:
        else{return false; }} // if(this.age.intValue()<6)...15
```

```
- // This class is thread-safe.
public class Person {
    private Integer age; // Shared variable
    private ReentrantLock lock = new ReentrantLock();

    public void setAge(Integer age){
        lock.lock();
        this.age = age;
        lock.unlock(); }

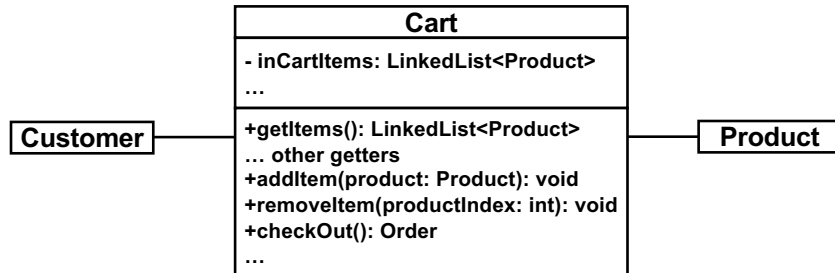

    public Integer getAge(){
        lock.lock();
        return this.age;
        lock.unlock(); }

    public boolean isKindergartener(){
        lock.lock();
        if(this.age < 6){return true;}
        else{return false;}
        lock.unlock(); }}
```



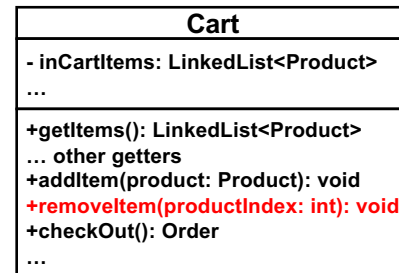


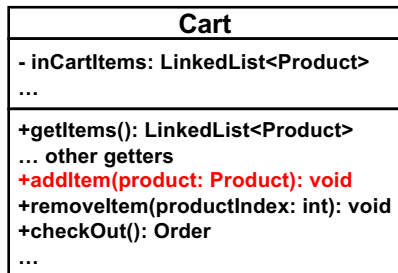


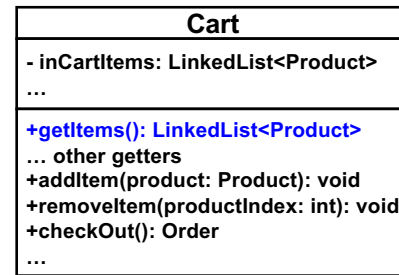
16

Exercise


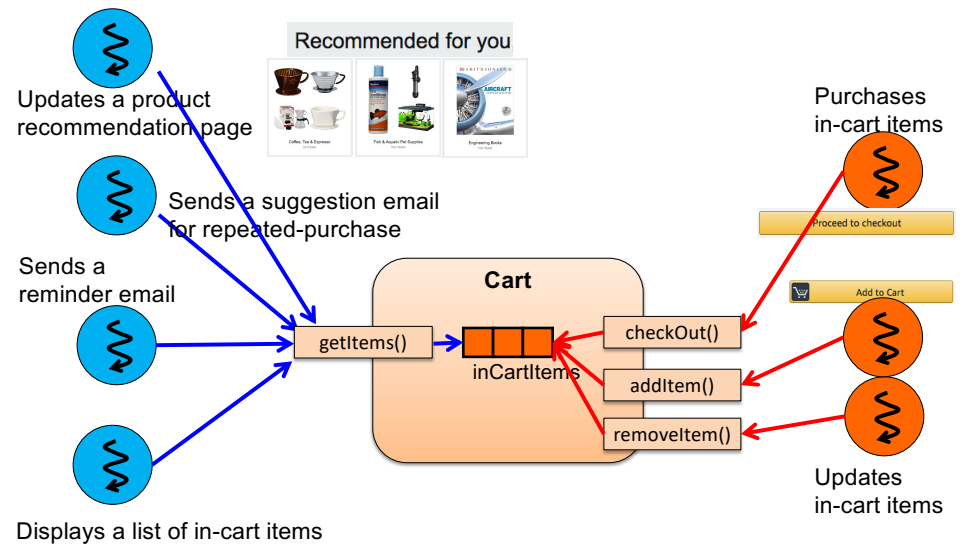


- Suppose you are implementing shopping carts for an online retailer.




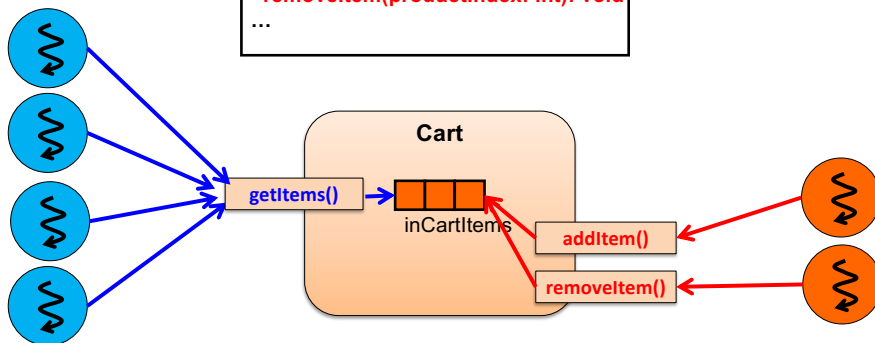
17



Cart
- inCartItems: LinkedList<Product> ...
+getItems(): LinkedList<Product> ... other getters +addItem(product: Product): void +removeItem(productIndex: int): void +checkout(): Order ...



Cart
- inCartItems: LinkedList<Product> ...
+getItems(): LinkedList<Product> +addItem(product: Product): void +removeItem(productIndex: int): void ...



```
class Cart{
    private LinkedList<Product> inCartItems
        = new LinkedList<>();


    public LinkedList<Product> getItems(){
        return inCartItems; } // READ

    public void addItem(Product item){
        inCartItems.add(item); } // WRITE

    public void removeItem(int productIndex){
        inCartItems.remove(productIndex); // WRITE
    } }
```

- **inCartItems** is shared by multiple threads.
- Need to guard it against concurrent access.
 - Identify every single *read* and *write* operation on it.
 - Surround it with `lock()` and `unlock()` on a lock

- **Cart is not thread-safe. Explain why and how to make it thread-safe.**
 - Note: `LinkedList` is not thread-safe. (All of its public methods such as `add()`, `remove()` and `get()` never perform thread synch.)




```
class Cart{
    private LinkedList<Product> inCartItems
        = new LinkedList<Product>();

    public LinkedList<Product> getItems(){
        return inCartItems; } // READ (2 steps)

    public void addItem(Product item){
        inCartItems.add(item); } // WRITE
                                // (multiple steps)
    public void removeItem(int productIndex){
        inCartItems.remove(productIndex); } // WRITE
                                                // (multiple steps)
} }
```

- Each write op is performed with a method of `LinkedList`.
 - `get()` and `remove()`
 - **Not thread-safe**
- If a thread calls `addItem()` and another calls `removeItem()` concurrently...
- If multiple threads call `removeItem()` concurrently...
- If multiple threads call `addItem()` concurrently...



```
class Cart{
    private LinkedList<Product> inCartItems
        = new LinkedList<Product>();

    public LinkedList<Product> getItems(){
        return inCartItems; } // READ (2 steps)

    public void addItem(Product item){
        inCartItems.add(item); } // WRITE
                                // (multiple steps)
    public void removeItem(int productIndex){
        inCartItems.remove(productIndex); } // WRITE
                                                // (multiple steps)
} }
```

- Each write op is performed with a method of `LinkedList`.
 - `get()` and `remove()`
 - **Not thread-safe**.
- We cannot change the methods of `LinkedList`.
 - Cannot introduce thread sync into those API methods.
- We should do thread sync in `Cart` (i.e. **client code of `LinkedList`**).

**“Case 2” where race conditions occur.
Do client-side locking.**

Thread-safe Shopping Cart


```
class Cart{
    private LinkedList<Product> inCartItems = new ...
                                // LinkedList is not thread-safe
    private ReentrantLock lock = new ...;

    public LinkedList<Product> getItems(){
        lock.lock();
        return inCartItems; // READ
        lock.unlock(); }

    public void addItem(Product item){
        lock.lock();
        inCartItems.add(item); // WRITE
        lock.unlock(); }

    public void removeItem(int productIndex){
        lock.lock();
        inCartItems.remove(productIndex); // WRITE
        lock.unlock(); } }
```

Another Exercise: Access Counter in a Concurrent Web Server



(Runnable class)


run(): Parses a request, reads a requested file, returns the file and **increments the file's access count**.

- Each request handler (thread):
 - Increments the access count of a file that is being accessed.


Concurrent Access Counter

- **AccessCounter**

- Contains a map that pairs a **file path** and its **access count**.
 - Assume `java.util.HashMap<java.nio.file.Path, Integer>`
- **void increment(Path path)**
 - accepts a file path and increments its access count.
 - if(**A requested path is in AC**){
 - increment the path's access count. }
 - else{
 - add the path and the access count of 1 to AC. }
- **int getCount(Path path)**
 - accepts a file path and returns its access count.
 - if(**A requested path is in AC**){
 - get the path's access count and return it. }
 - else{
 - return 0. }



- **HashMap** is **NOT thread-safe**.
 - Its public methods never perform thread synchronization.
 - `containsKey()`, `put()`, `get()`, `putIfAbsent()`, `replace()`, etc.
 - C.f. API doc: “[Note that this implementation is not synchronized.](#)”
 - Race conditions can occur in those public methods.
- **Client code of those public methods** needs to perform thread sync (client-side locking; c.f. **Case 2**).
 - **AccessCounter's increment() and getCount()**
 - **increment()**
 - if(**A requested path is in AC**){
 - increment the path's access count. }
 - else{
 - add the path and the access count of 1 to AC. }
 - **getCount()**
 - if(**A requested path is in AC**){
 - get the path's access count and return it. }
 - else{
 - return 0. }




HW 11

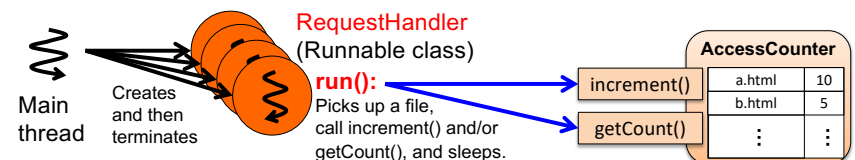
- **AccessCounter's increment() and getCount() need to perform thread sync.**

- **increment()**
 - **lock.lock();**
 - if(**A requested path is in AC**){
 - increment the path's access count. }
 - else{
 - add the path and the access count of 1 to AC. }
 - lock.unlock();**

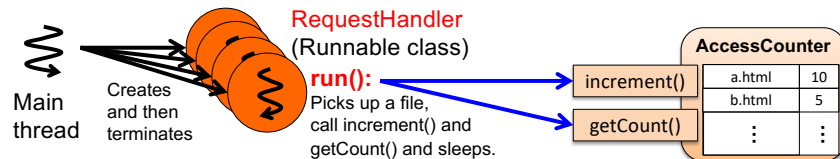
- **getCount()**
 - **lock.lock();**
 - if(**A requested path is in AC**){
 - get the path's access count and return it. }
 - else{
 - return 0. }
 - lock.unlock();**



- Implement **AccessCounter** as a **thread-safe Singleton** class.
 - Define a **HashMap<java.nio.file.Path, Integer>**
 - Define a regular (**non-static**) lock and use the lock to guard the **HashMap** in **increment()** and **getCount()**
 - Define another (**static**) lock and use the lock to guard the instance of **AccessCounter** in **getInstance()**
- Place some test/dummy files to be accessed (a.html, b.html, etc.)
- **RequestHandler (Runnable class)**
 - **run()**: Picks up one of the files at random, calls **increment()** and/or **getCount()** for that file, and sleeps for a few seconds. Repeats this forever with an infinite loop.
- Have **main()** create 10+ instances of **RequestHandler** and use 10+ threads to execute **RequestHandler's run()**.



- Do 2-step thread termination in `RequestHandler`.
 - Have the main thread terminate those 10+ extra threads in 2 steps.
 - Define a **volatile flag** in `RequestHandler`
 - Or, use an `AtomicBoolean`
 - Have the main thread flip the flag on each `RequestHandler` and call `interrupt()` on each request-handling threads.

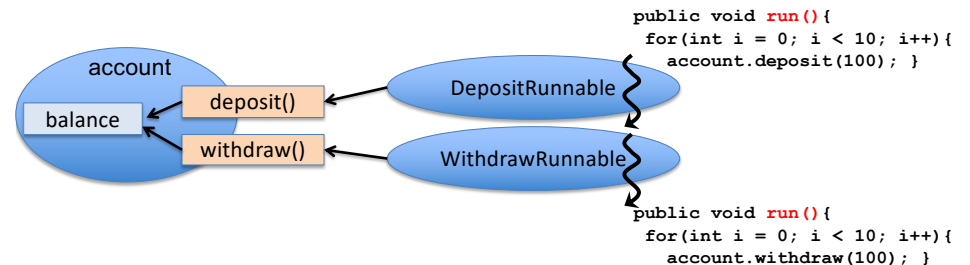


Thread Safety Issues

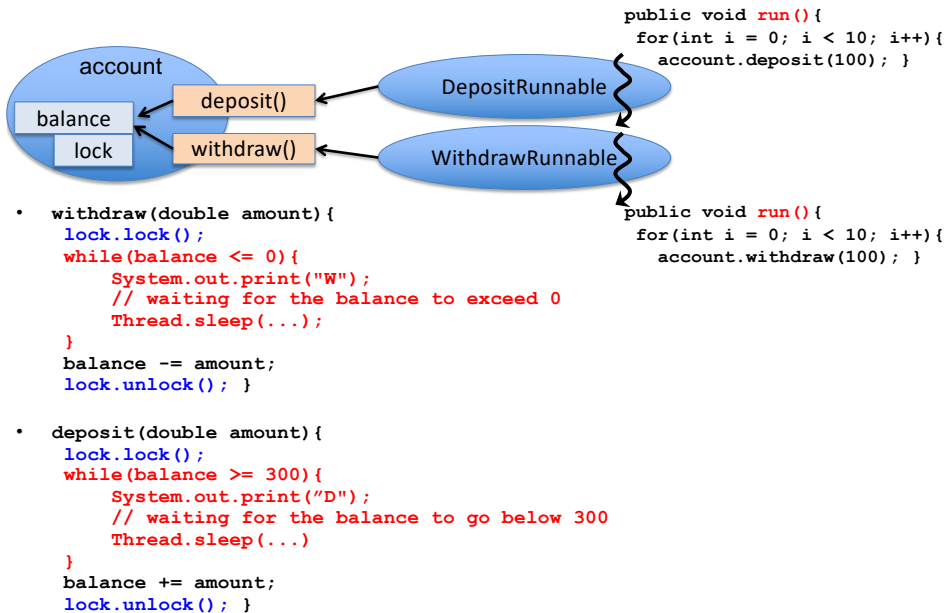
- Race conditions
 - Data inconsistency
- Deadlocks
 - Threads get stuck and cannot proceed even if they intend/want to do so.
- Thread-safe code is free from both race conditions and deadlocks.

Deadlock

DeadlockedBankAccount.java

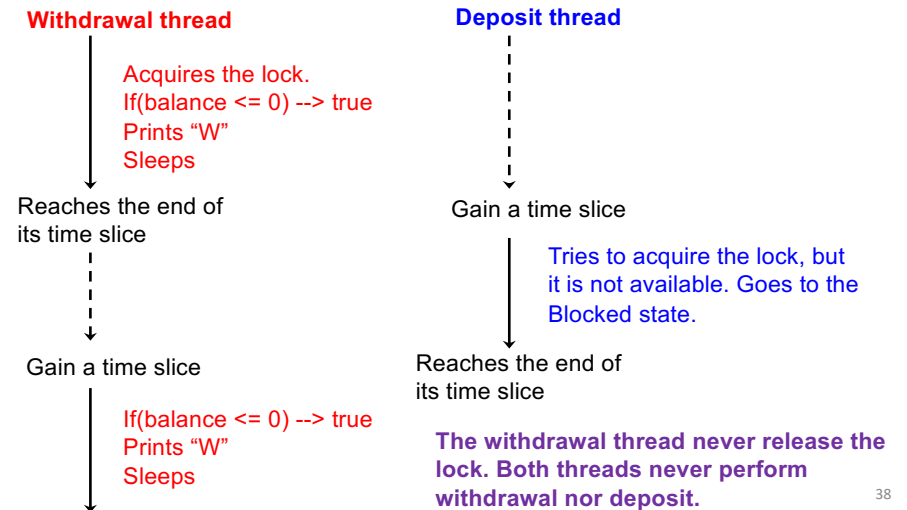


DeadlockedBankAccount.java



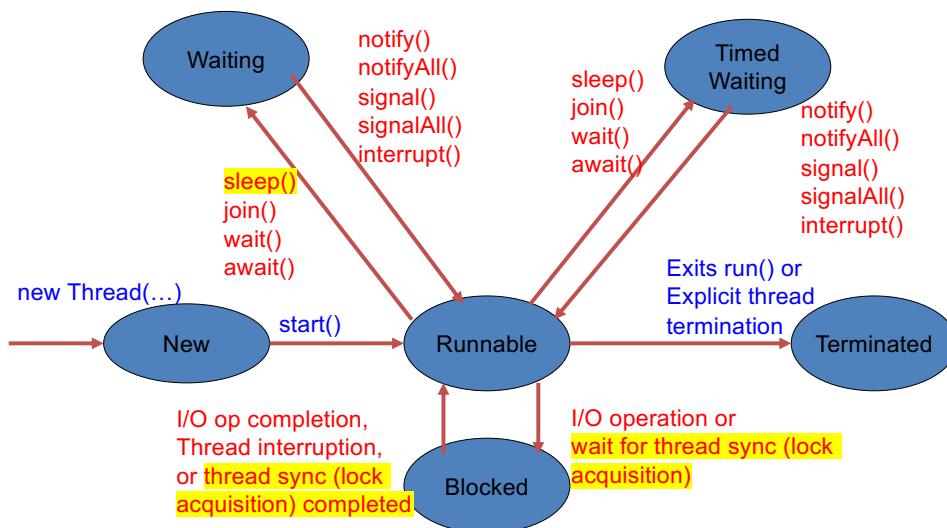
How Can a Deadlock Occur?

- Suppose the withdrawal thread goes ahead.



38

States of a Thread



39

Note

- A JVM can perform context switches even when a thread runs atomic code.
 - A lock guarantees that only one thread exclusively runs atomic code at a time.
 - It does NOT control when to (or when not to) perform context switches.
 - Some resources explicitly/implicitly say that context switches never occur when a thread runs atomic code.
 - It is WRONG!

DeadlockedBankAccount2.java

- Previous version

```
- withdraw(double amount){
    lock.lock();
    while( balance <= 0 ){
        System.out.print("W");
        Thread.sleep(...);
    }
    balance -= amount;
    lock.unlock();
}

- deposit(double amount){
    lock.lock();
    while( balance >= 300 ){
        System.out.print("W");
        Thread.sleep(...);
    }
    balance += amount;
    lock.unlock();
}
```

- New version

```
- withdraw(double amount){
    while( balance <= 0 ){
        System.out.print("W");
        Thread.sleep(...);
    }
    lock.lock();
    balance -= amount;
    lock.unlock();
}

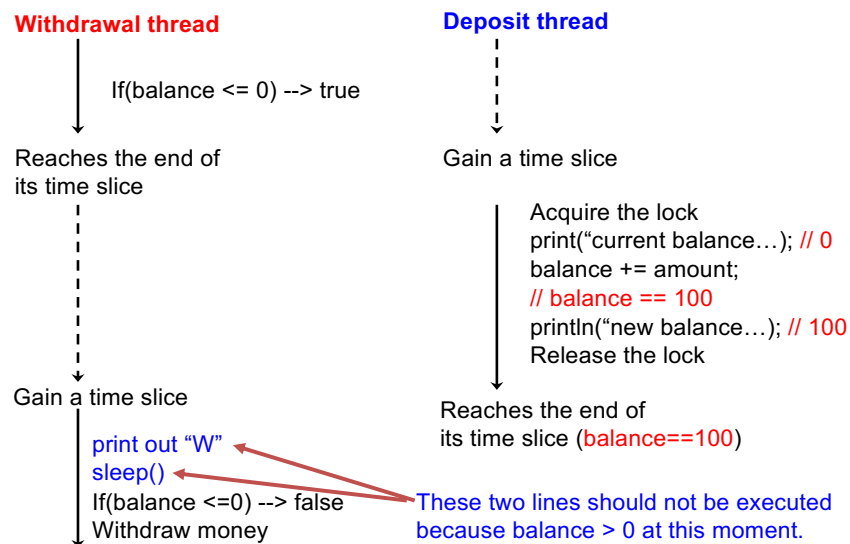
- deposit(double amount){
    while( balance >= 300 ){
        System.out.print("W");
        Thread.sleep(...);
    }
    lock.lock();
    balance += amount;
    lock.unlock();
}
```

41

- Has no deadlock problems.
- Can generate race conditions.

42

A Potential Race Condition in DeadlockedBankAccount2



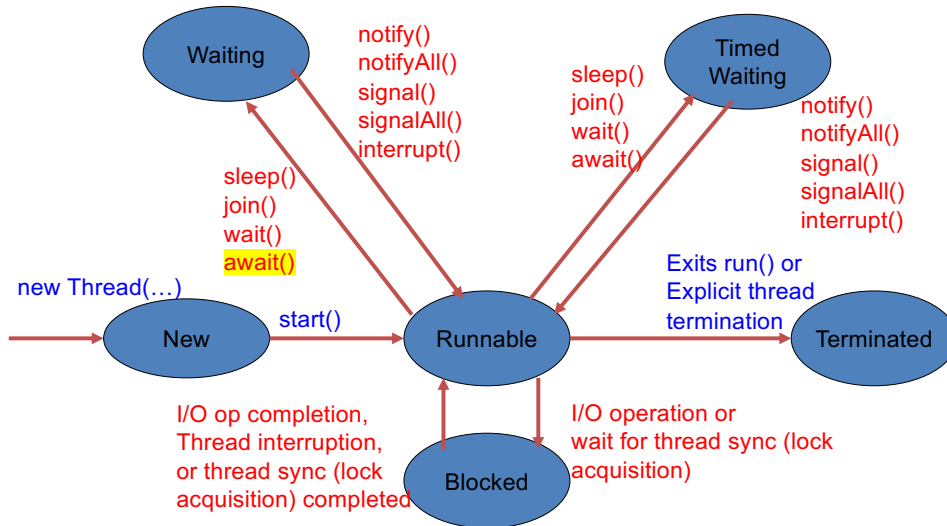
43

Avoiding Deadlocks and Race Conditions

- Use a **Condition** object.
 - java.util.concurrent.locks.Condition
 - Allows a thread to
 - Temporarily release a lock, so another thread can acquire it and proceed.
 - Re-acquire the lock later.
- Use a lock to obtain a condition object
 - ReentrantLock lock = new ReentrantLock();
 Condition condition = lock.newCondition(); // Factory method
 lock.lock();
 condition.await(); // Temporarily releases the lock.
 // Goes to the Waiting state until getting
 // signaled by another thread.

44

States of a Thread



45

ThreadSafeBankAccount2.java

- Condition `sufficientFundsCondition` = `lock.newCondition()`;
Condition `belowUpperLimitFundsCondition` = `lock.newCondition()`;
- `withdraw(double amount)`{
 `lock.lock()`;
 while(`balance <= 0`){
 // Wait for the balance to exceed 0
 `sufficientFundsCondition.await()`;
 }
 `balance -= amount`;
 `belowUpperLimitFundsCondition.signalAll()`;
 `lock.unlock()`;
}
- `deposit(double amount)`{
 `lock.lock()`;
 while(`balance >= 300`){
 // Wait for the balance to go below 300.
 `belowUpperLimitFundsCondition.await()`;
 }
 `balance += amount`;
 `sufficientFundsCondition.signalAll()`;
 `lock.unlock()`;
}

46

If a “Withdrawal” Thread Runs First and $balance \leq 0$...

- Condition `sufficientFundsCondition` = `lock.newCondition()`;
Condition `belowUpperLimitFundsCondition` = `lock.newCondition()`;
 - `withdraw(double amount)`{
 `lock.lock()`;
 while(`balance <= 0`){
 // Wait for the balance to exceed 0
 `sufficientFundsCondition.await()`;
 }
 `balance -= amount`;
 `belowUpperLimitFundsCondition.signalAll()`;
 `lock.unlock()`;
}
- // temporarily releases the lock and goes to the Waiting state til getting signaled by another thread

47

- Condition `sufficientFundsCondition` = `lock.newCondition()`;
Condition `belowUpperLimitFundsCondition` = `lock.newCondition()`;
 - `withdraw(double amount)`{
 `lock.lock()`;
 while(`balance <= 0`){
 // Wait for the balance to exceed 0
 `sufficientFundsCondition.await()`;
 }
 `balance -= amount`;
 `belowUpperLimitFundsCondition.signalAll()`;
 `lock.unlock()`;
}
 - `deposit(double amount)`{
 `lock.lock()`;
 while(`balance >= 300`){
 // Wait for the balance to go below 300.
 `belowUpperLimitFundsCondition.await()`;
 }
 `balance += amount`;
 `sufficientFundsCondition.signalAll()`;
 `lock.unlock()`;
}
- A “deposit” thread calls `signalAll()` to wake up a withdrawal thread(s) that is waiting until $balance > 0$.
It then releases the lock.

48

If a “Deposit” Thread Runs First and balance ≥ 300

```

• Condition sufficientFundsCondition = lock.newCondition();
Condition belowUpperLimitFundsCondition = lock.newCondition();

```

```

• withdraw(double amount){
    lock.lock();
    while(balance <= 0){
        // Wait for the balance to exceed 0
        sufficientFundsCondition.await(); }
    balance -= amount;
    belowUpperLimitFundsCondition.signalAll();
    lock.unlock(); }

```

The “withdrawal” thread
acquires the lock.

It then withdraws some
money.

```

• deposit(double amount){
    lock.lock();
    while(balance >= 300){
        // Wait for the balance to go below 300.
        belowUpperLimitFundsCondition.await(); }
    balance += amount;
    sufficientFundsCondition.signalAll();
    lock.unlock(); }

```

49

```

• Condition sufficientFundsCondition = lock.newCondition();
Condition belowUpperLimitFundsCondition = lock.newCondition();

```

```

• withdraw(double amount){
    lock.lock();
    while(balance <= 0){
        // Wait for the balance to exceed 0
        sufficientFundsCondition.await(); }
    balance -= amount;
    belowUpperLimitFundsCondition.signalAll();
    lock.unlock(); }

```

```

• deposit(double amount){
    lock.lock();
    while(balance >= 300){
        // Wait for the balance to go below 300.
        belowUpperLimitFundsCondition.await(); }
    balance += amount;
    sufficientFundsCondition.signalAll();
    lock.unlock(); }

```

Temporarily releases
the lock and goes to the
Waiting state until
balance < 300

50

```

• Condition sufficientFundsCondition = lock.newCondition();
Condition belowUpperLimitFundsCondition = lock.newCondition();

```

```

• withdraw(double amount){
    lock.lock();
    while(balance <= 0){
        // Wait for the balance to exceed 0
        sufficientFundsCondition.await(); }
    balance -= amount;
    belowUpperLimitFundsCondition.signalAll();
    lock.unlock(); }

```

A “withdrawal” thread
calls signalAll() to
wake up a “deposit”
thread(s) that are
waiting til balance
< 300.

It then releases the
lock.

```

• deposit(double amount){
    lock.lock();
    while(balance >= 300){
        // Wait for the balance to go below 300.
        belowUpperLimitFundsCondition.await(); }
    balance += amount;
    sufficientFundsCondition.signalAll();
    lock.unlock(); }

```

51

```

• Condition sufficientFundsCondition = lock.newCondition();
Condition belowUpperLimitFundsCondition = lock.newCondition();

```

```

• withdraw(double amount){
    lock.lock();
    while(balance <= 0){
        // Wait for the balance to exceed 0
        sufficientFundsCondition.await(); }
    balance -= amount;
    belowUpperLimitFundsCondition.signalAll();
    lock.unlock(); }

```

```

• deposit(double amount){
    lock.lock();
    while(balance >= 300){
        // Wait for the balance to go below 300.
        belowUpperLimitFundsCondition.await(); }
    balance += amount;
    sufficientFundsCondition.signalAll();
    lock.unlock(); }

```

The “withdrawal” thread
acquires the lock.

It then deposit some
money.

52

Two-Way (Bidirectional) Signaling b/w 2 Types of Threads

```

• Condition sufficientFundsCondition = lock.newCondition();
  Condition belowUpperLimitFundsCondition = lock.newCondition();

• withdraw(double amount){
  lock.lock();
  while(balance <= 0){
    // Wait for the balance to exceed 0
    sufficientFundsCondition.await();
  }
  balance -= amount;
  belowUpperLimitFundsCondition.signalAll();
  lock.unlock();
}

• deposit(double amount){
  lock.lock();
  while(balance >= 300){
    // Wait for the balance to go below 300.
    belowUpperLimitFundsCondition.await();
  }
  balance += amount;
  sufficientFundsCondition.signalAll();
  lock.unlock();
}

```


53

Condition

- **await()**
 - Blocks (i.e., gets stuck, or does not return) until getting signaled or interrupted
 - Blocks until getting signaled or interrupted, or until a specified waiting time (relative time) has elapsed.
 - Blocks until getting signaled or interrupted, or until a specified deadline (absolute time).
 - Throws an `InterruptedException`, if getting interrupted.
 - c.f. A previous lecture note on thread interruption
- **signalAll()**
 - Wakes up all the threads that wait on a condition object.
 - All of them go to the “Runnable” state and race to re-acquire a lock.
 - One of them will successfully re-acquire the lock.
 - All others will go to the “Blocked” state upon lock re-acquisition failures.
 - There are no ways to control which one of them to re-acquire the lock.

54


States of a Thread



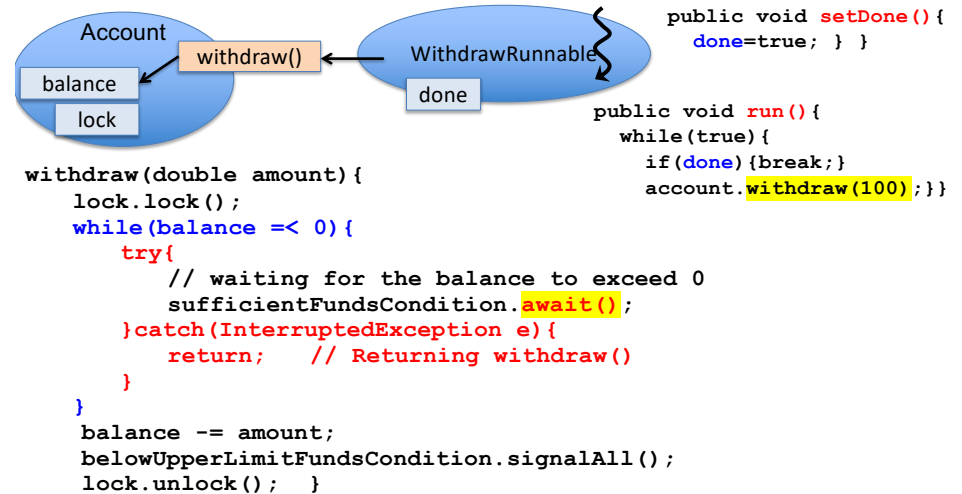
55

- When a thread calls **await()** or **signalAll()** on a **Condition** object,
 - the thread is assumed to hold a lock associated with the **Condition** object.
 - If the thread does not, an `IllegalMonitorStateException` is thrown.

Introducing 2-Step Thread Termination in ThreadSafeBankAccount2.java

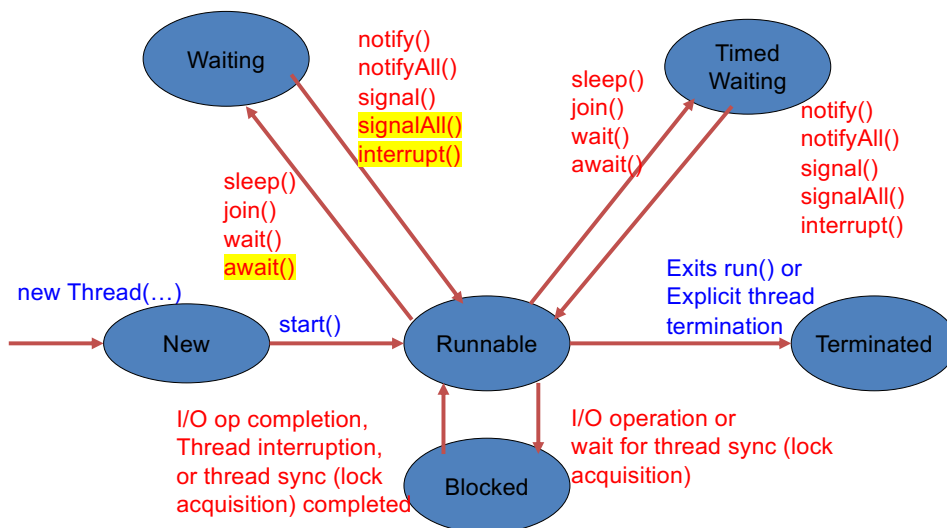


- Define a **volatile flag** (`done`) in `WithdrawRunnable`
 - Or, use an `AtomicBoolean`
- Have the main thread terminate each “withdrawal” thread by flipping the flag with `setDone()`.



- Have the main thread call `interrupt()`, after calling `setDone()`, on each “withdrawal” thread.
 - To let “withdrawal” threads wake up and catch an `InterruptedException`, if they are in the Waiting state due to `await()`.

States of a Thread



- If you don’t do 2-step thread termination but do a simple flag-based thread termination (i.e., if the main thread doesn’t call `interrupt()`)...
 - The main thread may not be able to terminate “withdrawal” threads forever.
 - If they are in the Waiting state and no other threads deposit sufficient money.
 - The bank app will never shut down.

HW 12

- Complete 2-step thread termination in `ThreadSafeBankAccount2.java`
 - in `WithdrawRunnable` and `DepositRunnable`
 - in `withdraw()` and `deposit()` threads.
- Run multiple “withdrawal” and “deposit” threads
- Have the main thread terminate both “withdrawal” and “deposit” threads.

```

• withdraw(double amount){
    lock.lock();
    while(balance <= 0){
        // waiting for the balance to exceed 0
        sufficientFundsCondition.await(); }
    belowUpperLimitFundsCondition.signalAll();
    balance -= amount;
    lock.unlock(); }

• deposit(double amount){
    lock.lock();
    while(balance >= 300){
        // waiting for the balance to go below 300.
        belowUpperLimitFundsCondition.await(); }
    sufficientFundsCondition.signalAll();
    balance += amount;
    lock.unlock(); }

```

- For example, do you have to worry about potential race conditions in this case? The answer is NO.

`signalAll()` Before or After a State Change?

```

• withdraw(double amount){
    lock.lock();
    while(balance <= 0){
        // waiting for the balance to exceed 0
        sufficientFundsCondition.await(); }
    balance -= amount;
    belowUpperLimitFundsCondition.signalAll();
    lock.unlock(); }

• deposit(double amount){
    lock.lock();
    while(balance >= 300){
        // waiting for the balance to go below 300.
        belowUpperLimitFundsCondition.await(); }
    balance += amount;
    sufficientFundsCondition.signalAll();
    lock.unlock(); }

```

- What if you call `signalAll()` first and then update the balance? Does it cause any problems?

Suppose `balance==0`.

(1) W thread:

```

Calls await().
Releases the
lock temporarily
and goes to the
“waiting” state.
Then, assume
a context switch
here.

```

```

withdraw(double amount){
    lock.lock();
    while(balance <= 0){
        // waiting for the balance to exceed 0
        sufficientFundsCondition.await(); }
    belowUpperLimitFundsCondition.signalAll();
    balance -= amount;
    lock.unlock(); }

```

(2) D thread:

```

Calls signalAll().
Then, assume a
context switch.

```

```

deposit(double amount){
    lock.lock();
    while(balance >= 300){
        // waiting for the balance to go below 300.
        belowUpperLimitFundsCondition.await(); }
    sufficientFundsCondition.signalAll();
    balance += amount;
    lock.unlock(); }

```

D thread: wakes up W thread, which is waiting for the balance to be >0

Can the “W” thread withdraw money before the “D” thread deposits money? (If this is possible, it’s a problem.) The answer is NO.

(1) W thread:
Calls `await()`.
Releases the lock temporarily and goes to the "waiting" state. Then, assume a context switch here.

```
withdraw(double amount){
    lock.lock();
    while(balance <= 0){
        // waiting for the balance to exceed 0
        sufficientFundsCondition.await();
    }
    belowUpperLimitFundsCondition.signalAll();
    balance -= amount;
    lock.unlock();
}
```

(3) W thread:
Wakes up and goes to the "runnable" state. Tries to acquire the lock again and **fails**. Goes to the "blocked" state.

(2) D thread:
Calls `signalAll()`. Then, assume a context switch.

```
deposit(double amount){
    lock.lock();
    while(balance >= 300){
        // waiting for the balance to go below 300.
        belowUpperLimitFundsCondition.await();
    }
    sufficientFundsCondition.signalAll();
    balance += amount;
    lock.unlock();
}
```

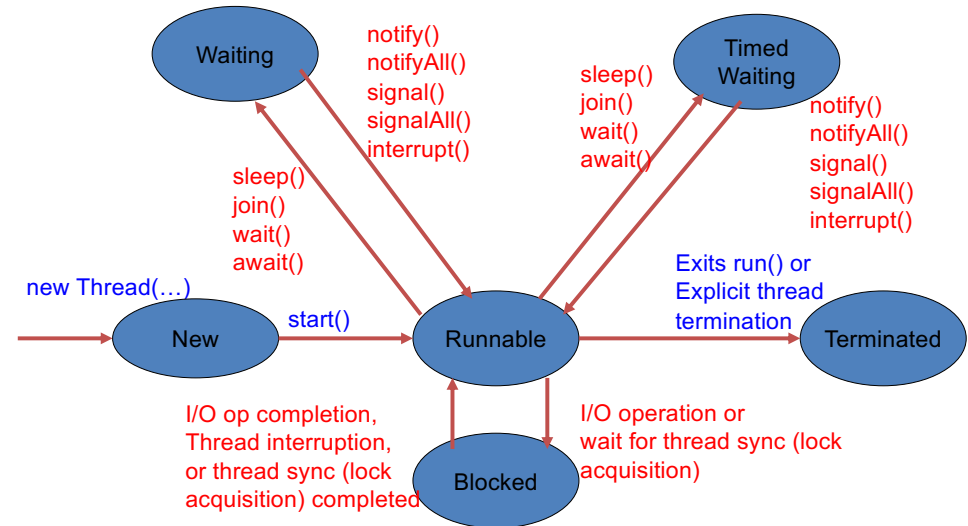
D thread: wakes up W thread

- "W" thread CANNOT withdraw money before "D" thread deposits money.

Two Important Things (1)

- You can safely change the value of a shared variable even after calling `signalAll()`.
 - AS FAR AS the value changes in atomic code (i.e. before calling `unlock()`)
- That said, common practice is:
 - A state change first, followed by `signalAll()`.

States of a Thread



66

Two Important Things (2)

- A JVM can perform context switches even when a thread runs atomic code.
 - A lock guarantees that only one thread exclusively runs atomic code at a time.
 - It does NOT control when to (or when not to) perform context switches.
 - Some resources explicitly/implicitly say that context switches never occur when a thread runs atomic code.
 - It is WRONG!

signal() and signalAll()

- `signalAll()`
 - Wakes up all “waiting” threads on a condition object.
 - All of them go to the “runnable” state.
 - They race to re-acquire a lock. Only one of them will actually re-acquire it. The others will go to the “blocked” state.
- `signal()`
 - Wakes up one of “waiting” threads on a condition object.
 - The selected thread goes to the “runnable” state. The others stay at the “waiting” state.
 - JVM’s thread scheduler selects one of them. Assume a random selection.
 - Not predictable which waiting thread to be selected.

“while” or “if” to Surround await()?

- ```
withdraw(double amount){
 lock.lock();
 while(balance <= 0){
 // waiting for the balance to exceed 0
 sufficientFundsCondition.await();
 }
 balance -= amount;
 belowUpperLimitFundsCondition.signalAll();
 lock.unlock(); }
```
- ```
deposit(double amount){
    lock.lock();
    while(balance >= 300){
        // waiting for the balance to go below 300.
        belowUpperLimitFundsCondition.await();
    }
    balance += amount;
    sufficientFundsCondition.signalAll();
    lock.unlock(); }
```
- “while” should be used rather than “if” when multiple threads call `withdraw()`/`deposit()` concurrently. Why?

signal() and signalAll()?

- Either one works well.
- `signalAll()` is favored in many cases/projects.
 - I prefer `signalAll()` in my personal taste.

A Potential Problem with “if”

(1) Suppose `balance==0`.
Two W threads call `await()`. They go to the “waiting” state and release the lock temporarily.
Then, assume a context switch.

(3) Two W threads go to the “runnable” state.
One of them acquires the lock again. The other one goes to the “blocked” state.

The 1st W thread withdraws \$100 (`balance==0`).

(2) D thread:
Deposits \$100.
Calls `signalAll()`.
Then, assume a context switch.

```
withdraw(double amount){
    lock.lock();
    if(balance <= 0){
        // waiting for ...
        sufficientFundsCondition.await();
    }
    balance -= amount;
    belowUpperLimitFundsCondition.signalAll();
    lock.unlock(); }

deposit(double amount){
    lock.lock();
    if(balance >= 300){
        // waiting for the balance to go below 300.
        belowUpperLimitFundsCondition.await();
    }
    balance += amount;
    sufficientFundsCondition.signalAll();
    lock.unlock(); }
```

(1) Suppose `balance==0`.

Two W threads call `await()`. They go to the "waiting" state and release the lock temporarily. Then, assume a context switch.

```
withdraw(double amount){
    lock.lock();
    if(balance <= 0){
        // waiting for ...
        sufficientFundsCondition.await();
    }
    balance -= amount;
    belowUpperLimitFundsCondition.signalAll();
    lock.unlock(); }
```

(3) Two W threads go to the "runnable" state. One of them acquires the lock again. The other one goes to the "blocked" state.

The 1st W thread withdraws \$100 (`balance==0`).

(4) The 2nd W thread acquires the lock after the 1st one releases it, and then withdraws \$100 (`balance==100`).

(2) D thread: Deposits \$100. Calls `signalAll()`. Then, assume a context switch.

```
deposit(double amount){
    lock.lock();
    if(balance >= 300){
        // waiting for the balance to go below 300.
        belowUpperLimitFundsCondition.await(); }
    balance += amount;
    sufficientFundsCondition.signalAll();
    lock.unlock(); }
```

(1) Suppose `balance==0`.

Two W threads call `await()`. They go to the "waiting" state and release the lock temporarily. Then, assume a context switch.

```
withdraw(double amount){
    lock.lock();
    if(balance <= 0){
        // waiting for ...
        sufficientFundsCondition.await();
    }
    balance -= amount;
    belowUpperLimitFundsCondition.signalAll();
    lock.unlock(); }
```

(3) Two W threads go to the "runnable" state. One of them acquires the lock again. The other one goes to the "blocked" state.

The 1st W thread withdraws \$100 (`balance==0`).

(4) The 2nd W thread acquires the lock after the 1st one releases it, and then withdraws \$100 (`balance==100`).

(2) D thread: Deposits \$100. Calls `signalAll()`. Then, assume a context switch.

```
deposit(double amount){
    lock.lock();
    if(balance >= 300){
        // waiting for the balance to go below 300.
        belowUpperLimitFundsCondition.await(); }
    balance += amount;
    sufficientFundsCondition.signalAll();
    lock.unlock(); }
```

- The 2nd "W" thread **should have made sure** if `balance>0` in (4).
- If only one "W" thread runs, this problem does not occur.
- Just always use **a while loop** regardless of the number of threads you use.

"if" in Atomic Code?

- If you want, you can use "if", rather than "while," for the conditional in atomic code
 - if you use `signal()`, not `signalAll()`.
- However, in practice, the **while-signalAll** pair is more common than the **if-signal** pair.