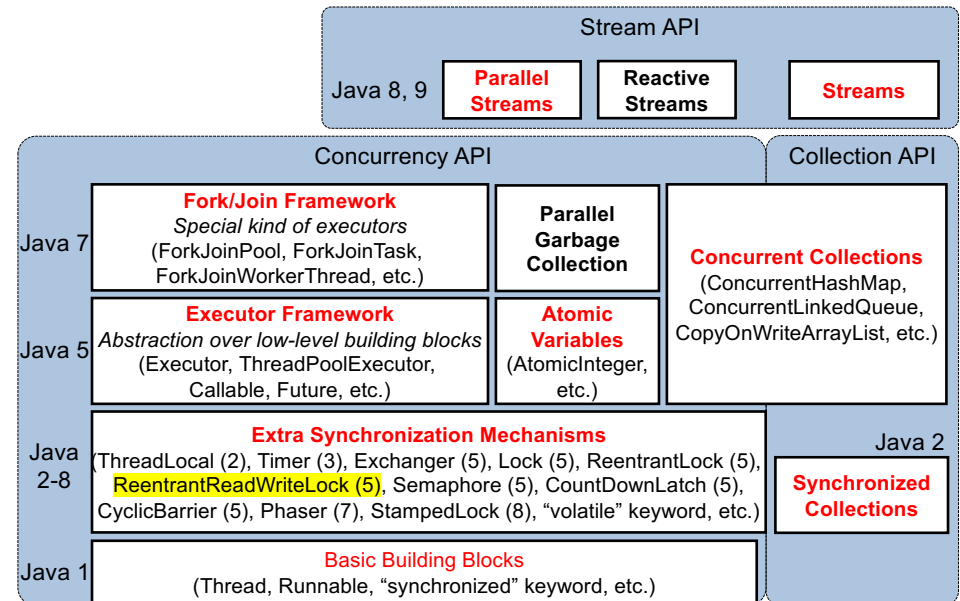


## Optimistic Thread Sync (Optimistic Locking) with Read-Write Locks

- Regular lock (`ReentrantLock`)
  - Used to avoid race conditions by mutually excluding multiple threads that access a shared variable.
    - Allows only one of them to access the variable at a time.
- Read-Write lock
  - A slight extension to `ReentrantLock`
  - A bit more *optimistic* than `ReentrantLock` to seek performance improvement.
    - `java.util.concurrent.locks.ReentrantReadWriteLock`

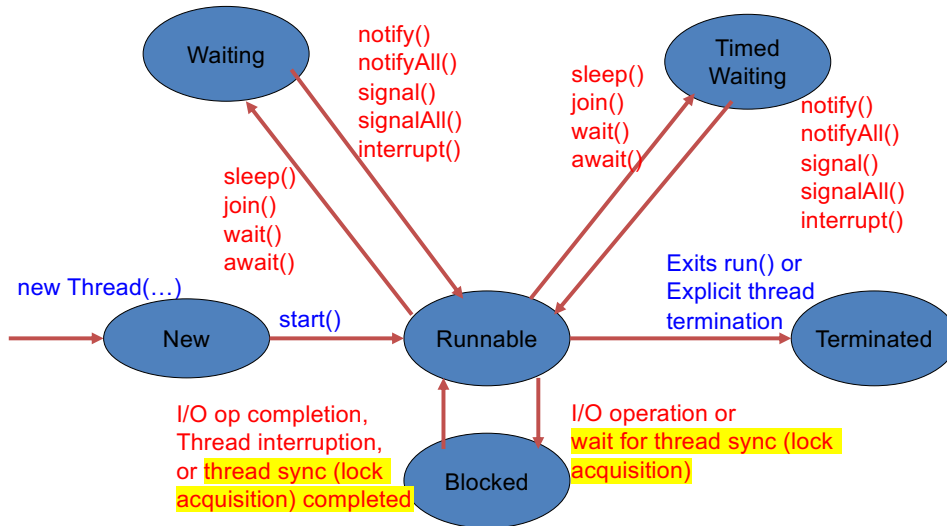
## Concurrency API in Java



## Thread Sync can be Computationally Expensive

- It takes some time for each thread to acquire and release a lock.
- If a lock is not available when a thread tries to acquire it, the thread is placed to the Blocked state.
  - It does nothing while it is in the Blocked state (i.e., until it can acquire the lock).
- If you have a lot of threads that compete for the same lock, many of them may not make any progress for a long time.


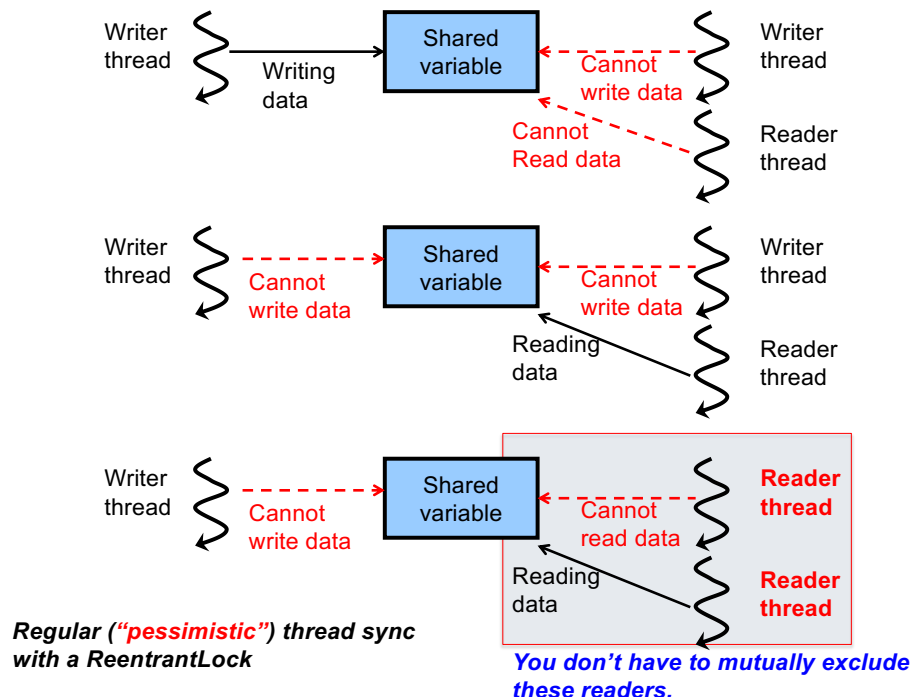
# States of a Thread




5

# Potential Performance Improvement

- If all threads are trying to read data from a shared variable, you **don't have to mutually exclude those "readers."**
  - The "reader" threads will never update the variable.
  - You never have to worry about race conditions if the value of a shared variable never change.
- You can be **optimistic** NOT to mutually exclude "reader" threads.



"Pessimistic" thread sync with a ReentrantLock



"Optimistic" thread sync with a ReentrantReadWriteLock

# ReentrantReadWriteLock

```
public class ReentrantReadWriteLock implements ReadWriteLock{  
    public class ReentrantReadWriteLock.ReadLock implements Lock{...}  
    public class ReentrantReadWriteLock.WriteLock implements Lock{...}  
    public ReentrantReadWriteLock.ReadLock readLock(){...}  
    public ReentrantReadWriteLock.WriteLock writeLock(){...} }
```

- Provides two locks as inner classes
  - **ReadLock** for reader threads to read data from a shared variable.
  - **WriteLock** for writer threads to write data to a shared variable.
- Provides factory methods for the two locks: `readLock()` and `writeLock()`.

9

- A reader can acquire a read lock even if it is already held by another reader,
  - **AS FAR AS** no writers hold a write lock.
- A writer can acquire a write lock **ONLY IF** no other writers and readers hold write/read locks.
  - Writers are mutually excluded as in pessimistic thread sync.

When another thread holds ...? Can a thread acquire...?	ReadLock	WriteLock
ReadLock	Yes	No
WriteLock	No	No

This turns to be NO if you use a regular "pessimistic" lock.

10

## When to Use Optimistic Thread Sync?

- When many reader threads run.
- When reader threads run more often than writer threads.
- When a read operation requires a long time to be completed.

## An Example Optimistic Thread Sync

```
int i; // shared variable  
ReentrantReadWriteLock rwLock = new ReentrantReadWriteLock();
```

- For reading data from the shared variable:

```
- rwLock.readLock().lock();  
try{  
    System.println(i); // atomic code  
    ...  
}finally{  
    rwLock.readLock().unlock(); }
```

- For writing data to the shared variable

```
- rwLock.writeLock().lock();  
try{  
    i++; // atomic code  
    ...  
}finally{  
    rwLock.writeLock().unlock(); }
```

11

# ReadLock and WriteLock

- Work similarly to `ReentrantLock`.
  - Support **nested locking** and **thread reentrancy**.
  - Support **interruption** via `Thread.interrupt()`.
- **WriteLock**
  - Returns a **Condition** object when `newCondition()` is called.
- **ReadLock**
  - Throws an `UnsupportedOperationException` when `newCondition()` is called.
  - Reader threads never need condition objects.
  - Reader threads never call `signalAll()` and `signal()`.


13

- **AccessCounter's** `increment()` and `getCount()` need to perform thread sync.

```

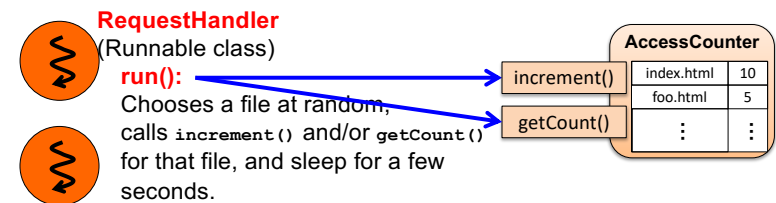
- increment()
    • lock.lock();
    if( A requested path is in AC ){
        increment the path's access count. }
    else{
        add the path and the access count of 1 to AC. }
    lock.unlock();

- getCount()
    • lock.lock();
    if( A requested path is in AC ){
        get the path's access count and return it. }
    else{
        return 0. }
    lock.unlock();
    
```



# HW 13

- Recall a previous HW to implement a concurrent access counter.
- **AccessCounter**
  - Maintains a map that pairs a **relative file path** and its **access count**.
    - Assume `java.util.HashMap<Path, Integer>`
  - void **increment**(Path path)
    - accepts a file path and increments its access count.
  - int **getCount**(Path path)
    - accepts a file path and returns its access count.




- Replace `ReentrantLock` with **ReentrantReadWriteLock** in `AccessCounter`

```

- increment()
    • rwLock.writeLock().lock();
    if( A requested path is in AC ){
        increment the path's access count. }           // Write
    else{
        add the path and the access count of 1 to AC. } // Write
    rwLock.writeLock().unlock();


- getCount()
    • rwLock.readLock().lock();
    if( A requested path is in AC ){
        get the path's access count and return it. }   // Read
    else{
        return 0. }
    rwLock.readLock().lock();
    
```



19

# Lock-ordering Deadlocks

## Lock-ordering Deadlocks



```

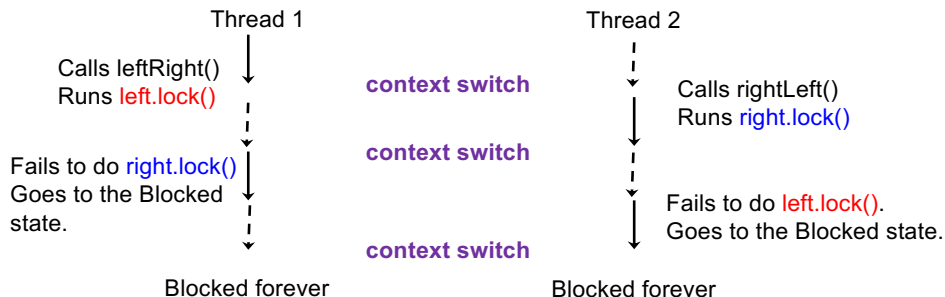
class LockOrderingDeadlock{
    private ReentrantLock left = new ReentrantLock();
    private ReentrantLock right = new ReentrantLock();

    public void leftRight(){
        left.lock();
        right.lock();
        // atomic code
        right.unlock();
        left.unlock();
    }

    public void rightLeft(){
        right.lock();
        left.lock();
        // atomic code
        left.unlock();
        right.unlock();
    }
}
    
```

A context switch can occur here.

A context switch can occur here.



```


class LockOrderingDeadlock{
    private ReentrantLock left = new ReentrantLock();
    private ReentrantLock right = new ReentrantLock();
    private ... ; // Shared variables

    public void leftRight(){
        left.lock();
        right.lock();
        // atomic code to access shared variables
        right.unlock();
        left.unlock();
    }

    public void rightLeft(){
        right.lock();
        left.lock();
        // atomic code to access shared variables
        left.unlock();
        right.unlock();
    }
}
    
```

- Problem:
  - Threads try to acquire *the same set of locks in different orders*.
    - Inconsistent lock ordering
      - Thread 1: left → right
      - Thread 2: right → left
- To-do:
  - Have all threads acquire the locks *in a globally-fixed order*.
- Be careful when you use multiple locks in order!**

# Dynamic Lock-ordering Deadlocks




```
class BankAccount{
    private double balance = 0;
    private ReentrantLock lock = ...;
    ...
    public void deposit(...)
    public void withdraw(...)
    public void transfer(...) }
```

```
void deposit(double amount){
    lock.lock();
    balance += amount;
    lock.unlock(); }
```

```
void withdraw(double amount){
    lock.lock();
    balance -= amount;
    lock.unlock(); }
```

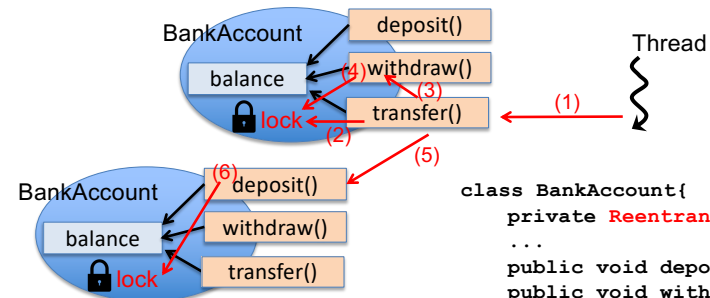
```
void void transfer(BankAccount destination, double amount){
    lock.lock();
    if( balance < amount )
        // Abort wire transfer. Generate an error msg or throw an exception
    else{
        withdraw(amount); // Nested locking. No problem.
        destination.deposit(amount); // Acquire another lock.
    }
    lock.unlock(); }
```



```
class BankAccount{
    private ReentrantLock lock = ...;
    ...
    public void deposit(...)
    public void withdraw(...)
    public void transfer(...) }
```

```
• void transfer(Account destination, double amount){
    lock.lock();
    if( balance < amount )
        // Abort wire transfer. Generate an error msg or throw an ...
    else{
        withdraw(amount); // Nested locking. No problem.
        destination.deposit(amount); // Acquire another lock.
    }
    lock.unlock(); }
```

- It looks as if all threads acquire the two locks (source account's lock and destination's lock) in the same order.
- However, this code can cause a lock-ordering deadlock.




```
class BankAccount{
    private ReentrantLock lock = ...;
    ...
    public void deposit(...)
    public void withdraw(...)
    public void transfer(...) }
```

```
void deposit(double amount){
    lock.lock();
    balance += amount;
    lock.unlock(); }
```

```
void withdraw(double amount){
    lock.lock();
    balance -= amount;
    lock.unlock(); }
```

```
void void transfer(BankAccount destination, double amount){
    lock.lock();
    if( balance < amount )
        // Abort wire transfer. Generate an error msg or throw an exception
    else{
        withdraw(amount); // Nested locking. No problem.
        destination.deposit(amount); // Acquire another lock.
    }
    lock.unlock(); }
```



Thread #1

Thread #2

Threads can deadlock in transfer().


- Threads #1 and #2 can deadlock in transfer() if
  - #1 transfers money from Account A to B
    - Acquires Account A's lock and B's lock.
  - #2 transfers money from B to A.
    - Acquires Account B's lock and A's lock.

```

• public void transfer(Account destination, double amount){
    lock.lock();
    if( balance < amount )
        // generate an error msg or...
    else{
        withdraw(amount);
        destination.deposit(amount);
    }
    lock.unlock();
}

```

A context switch can occur here.



## Solutions

- Static lock
- Timed locking
- Ordered locking
- Nested tryLock()

## Problem

– Threads try to acquire *the same set of locks in different orders*.

- Inconsistent lock ordering.

- Thread #1: Acct A's lock → Acct B's lock
- Thread #2: Acct B's lock → Acct A's lock

– This can occur with bad timing although code looks OK.

- A → B and C → D at the same time (No lock-ordering deadlock)
- A → B and A → C at the same time (No lock-ordering deadlock)
- A → B and B → A at the same time (Possible lock-ordering deadlock)

• **Be careful when you use multiple locks in order!!!**

## Solution 1: Static Lock

```

• private static ReentrantLock lock = new ReentrantLock();

• void deposit(double amount){
    lock.lock();
    balance += amount;
    lock.unlock();
}

void withdraw(double amount){
    lock.lock();
    balance -= amount;
    lock.unlock();
}

• public void transfer(Account destination, double amount){
    lock.lock();
    if( this.balance < amount )
        // generate an error msg or throw an exception
    else{
        this.withdraw(amount); // Nested locking
        destination.deposit(amount); // Nested locking!!!
    }
    lock.unlock(); // Make sure to release this lock when
                  // you abort wire transfer.
}

```

## Solution 2: Timed Locking

- Pros
  - Very simple
    - Uses only one lock (not two or more)
- Cons
  - Lack of concurrency
    - Deposit, withdrawal and transfer operations on different accounts are performed **sequentially** (not concurrently).
  - Performance penalty
    - All threads try to acquire a single lock. Higher chances that lock acquisition fails because the lock is not available.


```
• private ReentrantLock lock = new ReentrantLock();

• public void deposit(double amount){
    if( !lock.tryLock(3, TimeUnit.SECONDS) ){
        // generate an error msg or throw an exception
    }else{
        this.balance += amount;
        lock.unlock(); } }

• public void transfer(Account destination, double amount){
    lock.lock();
    if( this.balance < amount )
        // Abort wire transfer. Generate an error msg or ...
    else{
        this.withdraw(amount);    // Nested locking
        destination.deposit(amount);
    }
    lock.unlock(); }    // Make sure to release this lock when
                        // you abort wire transfer.
```

## Solution 3: Ordered Locking

- Pros
  - Simple
  - More efficient than Solution #1
    - By using a non-static lock
- Cons
  - Possibly unprofessional
    - May need to show an unprofessional message to the user
  - Transfers and deposits might not be completed in the worst case scenario.



- Define **a globally-fixed order** for `WireTransferHandler` to acquire two locks on two accounts
- Enforce the order in `WireTransferHandler.transfer()`



- An example of *globally-fixed* order
  - First, acquire the lock of an account with a **lower account #**
  - Then, acquire the lock of an account with a **higher acct #**

```

• public void transfer( Account source,
                      Account destination,
                      double amount){

    if( source.getAcctNum() < destination.getAcctNum() ){
        source.getLock().lock();
        destination.getLock().lock();
        if( source.getBalance() < amount )
            // generate an error msg or throw an exception
        else{
            source.withdraw(amount);    // Nested locking
            destination.deposit(amount); // Nested locking
        }
        destination.getLock().unlock();
        source.getLock().unlock();
    }
    else if( source.getAcctNum() > destination.getAcctNum() ){
        destination.getLock().lock();
        source.getLock().lock();
        ...
        source.getLock().unlock();
        destination.getLock().unlock(); } }

```

## Solution 3a: Ordered Locking with Instance IDs

- Instance IDs
  - Unique IDs (hash code) that the local JVM assigns to individual class instances.
    - Unique and intact on the same JVM
      - 2 instances of the same class have different IDs.
      - No instances share the same ID.
      - IDs never change after they are assigned to instances.
  - Use `System.identityHashCode(Object object)` OR `object.hashCode()`

- Pros
  - Locks are always acquired in a globally-fixed order.
  - More efficient than Solution #1
    - By using a non-static lock
  - More professional than Solution #2
    - Transfers and deposits can complete for sure.
- Cons
  - Rely on an application-specific/dependent data (acct #)
  - Once an account is set up, its account number should not be changed.
    - If you allow dynamic changes of an account number, you need to use an extra lock in `BankAccount` to guard it.

- An example of *globally-fixed* order
  - First, acquire the lock of an account with a **“smaller” ID**
  - Then, acquire the lock of an account with a **“bigger” ID**

```

• public void transfer( Account source,
                      Account destination,
                      double amount){

    int sourceID = source.hashCode();
    int destID = destination.hashCode();

    if( sourceID < destID ){
        source.getLock().lock();
        destination.getLock().lock();
        if( source.getBalance() < amount )
            // generate an error msg or throw an exception
        else{
            source.withdraw(amount);    // Nested locking
            destination.deposit(amount); // Nested locking
        }
        destination.getLock().unlock();
        source.getLock().unlock();
    }
    if( sourceID > destID ){
        destination.getLock().lock();
        source.getLock().lock();
        ...
        source.getLock().unlock();
        destination.getLock().unlock(); } }

```

## Solution 4: Nested Timed Locking

- Pros
  - Locks are always acquired in a globally-fixed order.
  - More efficient than Solution #1
    - By using a non-static lock
  - More professional than Solution #2
    - Transfers and deposits complete for sure.
  - No application-specific data (e.g., account numbers) are required to order locking.


- Cons
  - N/A

```

public void transfer(Account source,
                    Account destination,
                    double amount){
    Random random = new Random();
    while(true){
        if( source.getLock().tryLock() ){
            try{
                if( destination.getLock().tryLock() ){
                    try{
                        if( source.getBalance() < amount )
                            // generate an error msg/exception
                        else{
                            source.withdraw(amount);
                            destination.deposit(amount);
                        }
                        return;
                    }finally{
                        destination.getLock().unlock();
                    }
                }
            }finally{
                source.getLock().unlock();
            }
        }
        Thread.sleep( Duration.ofMillis(random.nextInt(50)));
    }
}

```

- If the first `tryLock()` fails, sleep and try again.
- If the first `tryLock()` succeeds but the second one fails, unlock the first lock, sleep, and try again.




- Use nested `tryLock()` calls to implement an **ALL-OR-NOTHING** policy.
  - Acquire both of A's and B's locks, OR
  - Acquire none of them.
- Avoid a situation where a thread acquires one of the two locks and fails to acquire the other.

- Pros
  - More efficient than Solution #1
    - By using a non-static lock
  - More professional than Solution #2
    - Transfers and deposits complete for sure.
  - No application-specific data (e.g., account numbers) are required to order locking.
- Cons
  - Not that simple

# Exercise: Real-time Admission Tracking


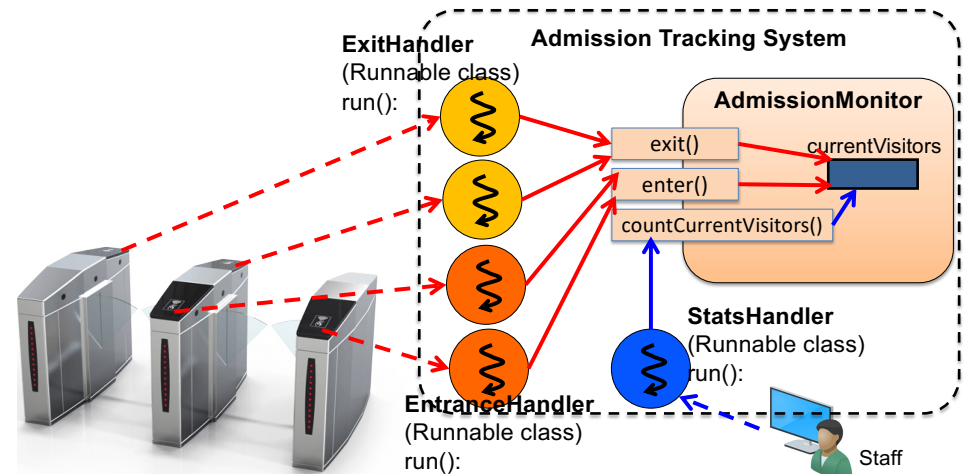
- Electronic, real-time admission tracking and control (in a museum, for example)



- Monitor the # of visitors who are currently in.
- Record the # of visitors per day.
- Record how long each visitor stays in.

- Monitor the number of visitors who are currently in.

AdmissionMonitor	
-	currentVisitors: int
+	countCurrentVisitors():int
+	enter(): void
+	exit(): void



```
class AdmissionMonitor{
    private int currentVisitors = 0;

    public void enter(){
        currentVisitors++;
    }

    public void exit(){
        currentVisitors--;
    }

    public int countCurrentVisitors(){
        return currentVisitors;
    }
}
```


```
class EntranceHandler implements ...{
    private AdmissionMonitor monitor;
    public void run(){
        monitor.enter();
    }
}
```

```
class ExitHandler implements ...{
    private AdmissionMonitor monitor;
    public void run(){
        monitor.exit();
    }
}
```

```
class StatsHandler implements ...{
    private AdmissionMonitor monitor;
    public void run(){
        monitor.countCurrentVisitors();
    }
}
```

AdmissionMonitor is not thread-safe.

## Conditional Admission in enter()



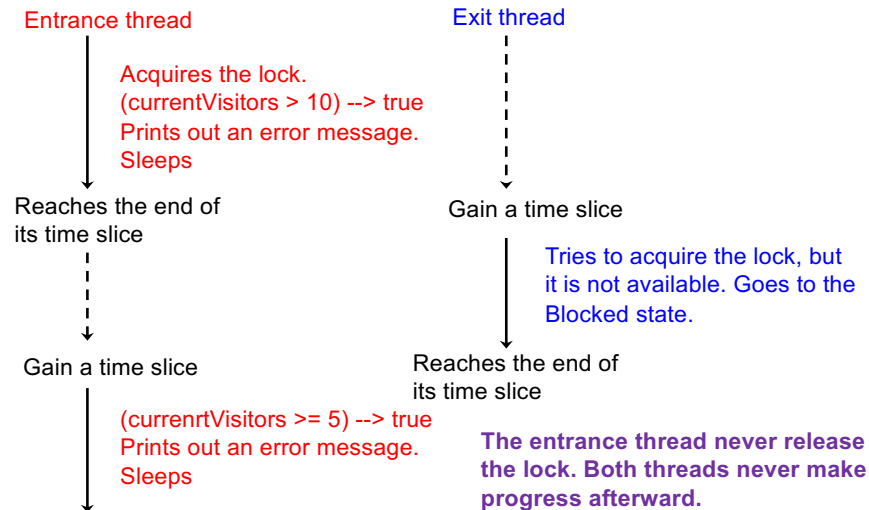
```
• enter(){
    lock.lock();
    while(currentVisitors > 10){
        System.out.println("Too many visitors. Please wait for a while!");
        // waiting until the # of visitors goes below 10.
        Thread.sleep(...);
    }
    currentVisitors++;
    lock.unlock();
}

• exit(double amount){
    lock.lock();
    currentVisitors--;
    lock.unlock();
}
```

This code causes a deadlock!

# How Can a Deadlock Occur?

- Suppose 15 visitors are already in.



49

## HW 13

- Submit a thread-safe version of `AdmissionMonitor`
  - Implement conditional admission
  - Avoid race conditions
    - `currentVisitors` is shared by 3 types of threads
  - Avoid deadlocks with a condition object
    - Define a `Condition` as a data field of `AdmissionMonitor`.
    - Replace `sleep()` with `await()`
    - Have `exit()` call `signalAll()`.
  - Use `ReentrantReadWriteLock` rather than `ReentrantLock`.
    - Use a write lock in `enter()` and `exit()`.
    - Use a read lock in `countCurrentVisitors()`.
- Implement 2-step thread termination for the main thread to terminate all “entrance”, “exit” and “stats” threads.

# Important Note

- DO NOT allow a thread to **conditionally stop making progress** (i.e., wait until a certain condition is satisfied) **with a lock held**.
  - Use a condition object, so the thread can temporarily release the lock.
  - c.f. Banking app example

## Exercise: New Feature in Shopping Carts

**Cart**

- `inCartItems: LinkedList<Product>`

- `subtotal: float`

---

`+getItems(): LinkedList<Product>`

`+getSubTotal(): float`

`+addItem(product: Product): void`

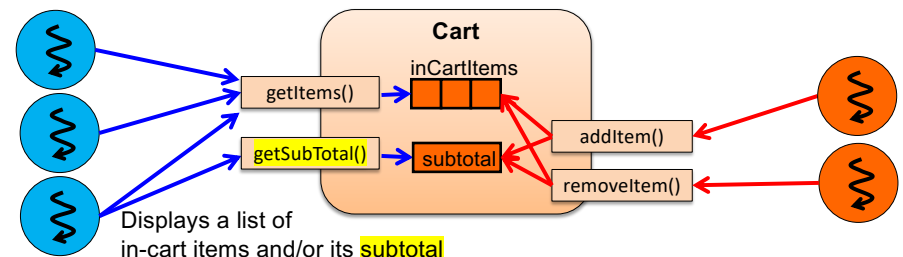
`+removeItem(productIndex: int): void`

...

Shopping Cart

	Price	Quantity
Uncommon Type by Tom Hanks Paperback In Stock ✓prime	\$11.00	1
<input type="checkbox"/> This is a gift Learn more <a href="#">Delete</a> <a href="#">Save for later</a> <a href="#">Compare with similar items</a>		
LEGO Technic Racing Yacht 42074 Building Kit (330 Pieces) In Stock ✓prime	\$31.99	1
<input type="checkbox"/> This is a gift Learn more <a href="#">Delete</a> <a href="#">Save for later</a>		

**Subtotal (2 items): \$42.99**



```

class Cart{
    private LinkedList<Product> inCartItems= new...
    private float subtotal;
    private ReentrantLock lock = new ...;


    public LinkedList<Product> getItems(){
        lock.lock();
        return inCartItems;
        lock.unlock();
    }

    public float getSubTotal(){
        return subtotal;
    }

    public void addItem(Product item){
        lock.lock();
        inCartItems.add(item);
        lock.unlock();
        subtotal += item.getPrice();
    }

    public void removeItem(int productIndex){
        subtotal -= inCartItems.get(productIndex).getPrice();
        lock.lock();
        inCartItems.remove(productIndex);
        lock.unlock();
    }
}

```



- Cart is not thread-safe. C.f. Case 1

## Using Solution 1 of Case 1

```

class Cart{
    private LinkedList<Product> inCartItems;
    private ReentrantLock lock = new ...;

    public LinkedList<Product> getItems(){
        lock.lock();
        return inCartItems;
        lock.unlock();
    }

    public void addItem(Product item){
        lock.lock();
        inCartItems.add(item);
        lock.unlock();
    }

    public void removeItem(int productIndex){
        lock.lock();
        inCartItems.remove(productIndex);
        lock.unlock();
    }

    public float getSubTotal(){
        float subtotal;
        lock.lock();
        for(Product item: inCartItems){
            subtotal += item.getPrice();
        }
        lock.unlock();
        return subtotal;
    }
}

```

- Turn a shared variable to be a local variable
  - Eliminate the data field subtotal.
  - Use a lock to guard inCartItems.

## Using Solution 2 in Case 1

```

class Cart{
    private LinkedList<Product> inCartItems;
    private float subtotal;
    private ReentrantLock lock = new ...;

    public LinkedList<Product> getItems(){
        lock.lock();
        return inCartItems;
        lock.unlock();
    }

    public float getSubTotal(){
        lock.lock();
        return subtotal;
        lock.unlock();
    }

    public void addItem(Product item){
        lock.lock();
        inCartItems.add(item);
        subtotal += item.getPrice();
        lock.unlock();
    }

    public void removeItem(int productIndex){
        lock.lock();
        subtotal -= inCartItems.get(productIndex).getPrice();
        inCartItems.remove(productIndex);
        lock.unlock();
    }
}

```

- Use a lock to guard inCartItems and subtotal.
- It is perfectly fine to guard multiple shared variables with a single lock.